



Database Systems

Lecture18 – Chapter 17: Transactions



Beomseok Nam (남범석)

bnam@skku.edu



Outline

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability.

Transaction Concept


- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
- E.g., transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- Two main issues to deal with:
 - Failures of various kinds, such as hardware failures and system crashes
 - Concurrent execution of multiple transactions

Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
 1. **read**(A)
 2. $A := A - 50$
 3. **write**(A)
 4. **read**(B)
 5. $B := B + 50$
 6. **write**(B)
- **Atomicity requirement**
 - If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - The system should ensure that updates of a partially executed transaction are not reflected in the database
- **Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.

Example of Fund Transfer (Cont.)

- **Consistency requirement** in above example:
 - The sum of A and B is unchanged by the execution of the transaction



- In general, consistency requirements include
 - Explicit integrity constraints such as primary keys and foreign keys
 - Implicit integrity constraints
 - e.g., sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
 - A transaction must see a consistent database.

Example of Fund Transfer (Cont.)

- **Isolation requirement:** if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

T1

1. **read**(A)
2. $A := A - 50$
3. **write**(A)
4. **read**(B)
5. $B := B + 50$
6. **write**(B)

T2

read(A), read(B), print($A+B$)

- Isolation can be ensured trivially by running transactions **serially**, i.e., one after the other.
- However, executing multiple transactions sequentially has performance problems.

ACID Principles


- To preserve the integrity of data the database system must ensure:
- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are. All or Nothing.
- **Consistency.** Atomic execution of a transaction in isolation preserves the consistency of the database.
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – state after the final statement has been executed.
- **Failed** – state after the discovery that normal execution can no longer proceed.
- **Aborted** – state after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.
 - Two options after it has been aborted:
 - Restart the transaction
 - Can be done only if no internal logical error
 - Kill the transaction
- **Committed** – state after successful completion.

Transaction State Transition Diagram





Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system.
Advantages are:
 - **Increased processor and disk utilization**, leading to better transaction *throughput*
 - E.g., one transaction can be using the CPU while another is reading from or writing to the disk
 - **Reduced average response time** for transactions: short transactions need not wait behind long ones.

- **Concurrency control schemes** – mechanisms to achieve isolation
 - control the interaction among the concurrent transactions
 - prevent from destroying the consistency of the database
 - Will study in Chapter 18



Schedules

- **Schedule** –specify the order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of all instructions of those transactions
 - Must preserve the order of instructions that appear in each individual transaction.

- A transaction that successfully completes its execution will have a commit instruction as the last statement

- A transaction that fails to complete its execution will have an abort instruction as the last statement

Serial Schedule 1

- Let T_1 transfer \$50 from A to B , and T_2 transfer 10% of the balance from A to B .
- A **serial** schedule in which T_1 is followed by T_2 :

| T_1 | T_2 |
|--|---|
| read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit | read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit |

Serial Schedule 2

- A serial schedule where T_2 is followed by T_1

| T_1 | T_2 |
|--|---|
| read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit | read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit |

Serializable Schedule 3

- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1

| T_1 | T_2 |
|--|---|
| read (A) $A := A - 50$ write (A) | |
| | read (A) $temp := A * 0.1$ $A := A - temp$ write (A) |
| read (B) $B := B + 50$ write (B) commit | |
| | read (B) $B := B + temp$ write (B) commit |

- In Schedules 1, 2 and 3, the sum $A + B$ is preserved.

Schedule 4

- The following schedule does not preserve the value of $(A + B)$.


| T_1 | T_2 |
|---|---|
| read (A) $A := A - 50$ | |
| | read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) |
| write (A) read (B) $B := B + 50$ write (B) commit | |
| | $B := B + temp$ write (B) commit |

Serializability


- **Basic Assumption:**

- We assume each transaction preserves database consistency.

- Thus, serial execution of a set of transactions preserves database consistency.



- A schedule is **serializable** if it is equivalent to a serial schedule.



- Different forms of serializability:

1. **Conflict serializability**
2. **View serializability**



Simplified view of transactions

- We ignore operations other than **read** and **write** instructions
- We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- Our simplified schedules consist of only **read** and **write** instructions.

Conflicting Instructions

- Instructions I_i and I_j of transactions T_i and T_j **conflict** if and only if at least one of these instructions write a shared data item Q .
 - $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 - $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 - $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict
 - $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict
- Note: a conflict forces an order between them.
- If *instructions* are consecutive in a schedule and they do not conflict, their results would remain the same even if we interchange them in the schedule.



Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**.
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule

Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore Schedule 3 is **conflict serializable**.

| T_1 | T_2 |
|-----------------------|-----------------------|
| read (A) write (A) | |
| | read (A) write (A) |
| read (B) write (B) | |
| | read (B) write (B) |

Schedule 3

| T_1 | T_2 |
|--|--|
| read (A) write (A) read (B) write (B) | |
| | read (A) write (A) read (B) write (B) |

Schedule 6

Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

| T_3 | T_4 |
|---------------|---------------|
| read (Q) | write (Q) |
| write (Q) | |

- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

View Serializability

- Let S and S' be two schedules with the same set of transactions.
- S and S' are **view equivalent** if the following three conditions are met
 1. If in schedule S , transaction T_i reads the initial value of Q , then in schedule S' also transaction T_i must read the initial value of Q .
 2. If in schedule S transaction T_i executes **read**(Q), and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same **write**(Q) operation of transaction T_j .
 3. The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S' .
- View equivalence is also based purely on **reads** and **writes** alone.

View Serializability (Cont.)

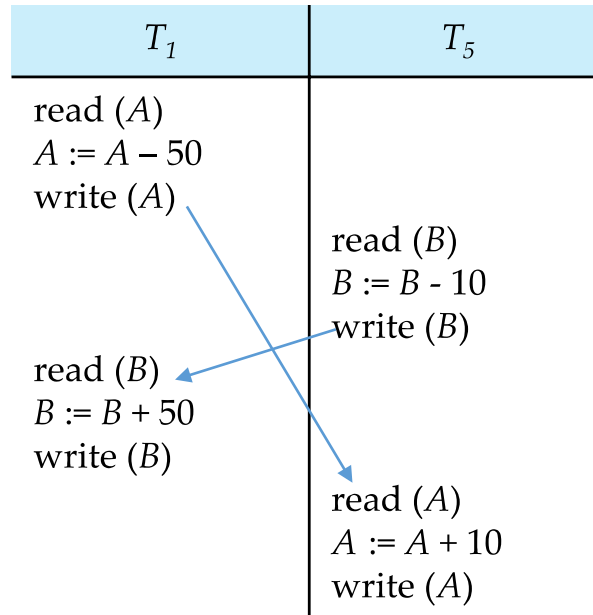
- A schedule S is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view-serializable but *not* conflict serializable.

| T_{27} | T_{28} | T_{29} |
|---------------|---------------|---------------|
| read (Q) | write (Q) | |
| write (Q) | | write (Q) |

- What serial schedule is above equivalent to?
- Every view serializable schedule that is not conflict serializable has **blind writes**.

Other Notions of Serializability


- The schedule below produces same outcome as the serial schedule $\langle T_1, T_5 \rangle$, yet is not conflict equivalent or view equivalent to it.



- Determining such equivalence requires analysis of operations other than read and write.




Testing for Serializability

- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** — a direct graph where the vertices are the transactions (names).
- We draw an arc from T_i to T_j if the two transactions conflict, and T_i accessed the conflicting data item earlier than T_j .
- We may label the arc by the item that was accessed.
- Example of a precedence graph



Test for Conflict Serializability

- A schedule is **conflict serializable** if and only if its precedence graph is **acyclic**.
- Cycle-detection algorithms exist which take order n^2 time, where n is the number of vertices in the graph.
 - Better algorithms take order $n + e$ where e is the number of edges.
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
 - This is a linear order consistent with the partial order of the graph.



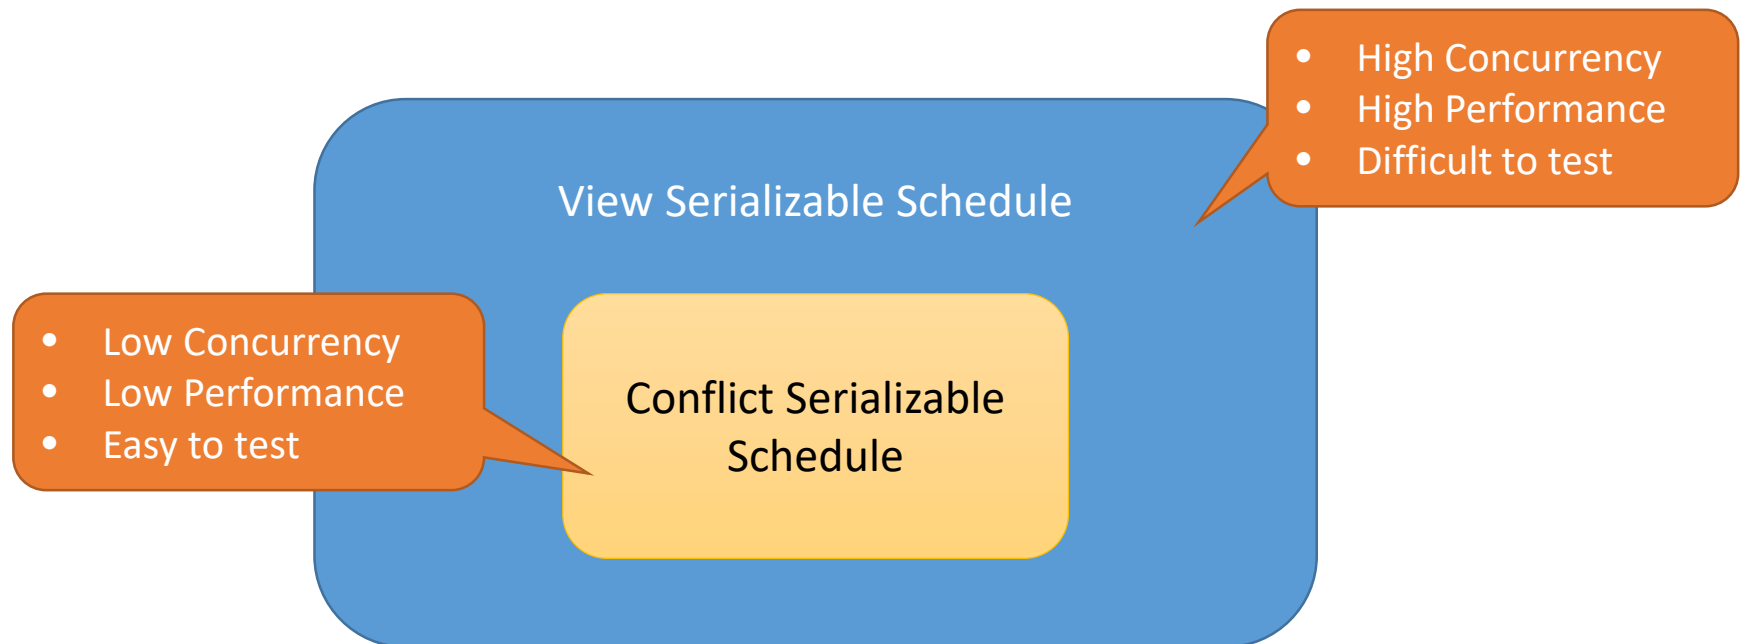


Test for View Serializability

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability.
 - Extension to test for view serializability has cost exponential in the size of the precedence graph.
- The problem of checking if a schedule is view serializable falls in the class of *NP*-complete problems.
 - Thus, existence of an efficient algorithm is *extremely* unlikely.
- However practical algorithms that just check some **sufficient conditions** for view serializability can still be used.

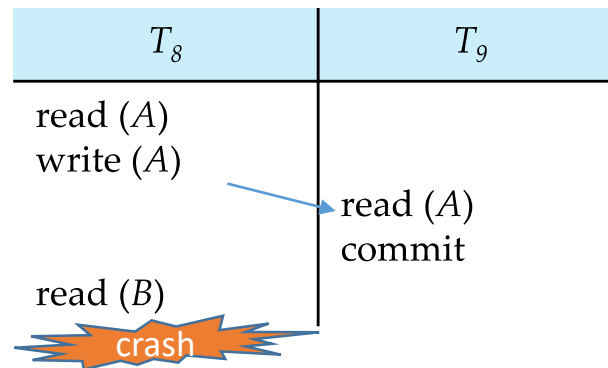
Conflict Serializability vs. View Serializability

- Conflict serializability is more popular in DBMS in practice
- But, conflict serializability actually permit only a subset of serializable schedules that do not have consistency problems.
- The general form of view serializability is very expensive to test, and only a very restricted form of it is used for concurrency control.



Recoverable Schedules

- Need to address the effect of transaction failures on concurrently running transactions.
- **Recoverable schedule** — if a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i appears before the commit operation of T_j .
- The following schedule is not recoverable



- If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

Cascading Rollbacks

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

| T_{10} | T_{11} | T_{12} |
|-----------------------------------|-----------------------|----------|
| read (A) read (B) write (A) | read (A) write (A) | read (A) |
| abort | | |

If T_{10} fails, T_{11} and T_{12} must also be rolled back.

- Can lead to the undoing of a significant amount of work



Cascadeless Schedules

- **Cascadeless schedules** — cascading rollbacks cannot occur;
 - For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
 - Read committed data only!
- Every Cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless



Concurrency Control

- **Goal:** A DBMS must make scheduling decisions that are
 - either conflict or view serializable, and
 - recoverable and preferably cascadeless

- DBMS scheduler should provide a high degree of concurrency
 - A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency.
 - Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur.
 - Some schemes allow only conflict-serializable schedules to be generated, while others allow view-serializable schedules that are not conflict-serializable.
 - We study concurrency control protocols in Chapter 18.



Weak Levels of Consistency

- Some applications (OLAP) are willing to live with weak levels of consistency, allowing schedules that are not serializable
 - Tradeoff accuracy for performance
- On Line Transaction Processing –OLTP
 - Maintains a database that is an accurate model of some real-world enterprise
 - Short simple transactions
 - Transactions access only a small fraction of the database
- On Line Analytic Processing –OLAP
 - Uses database to guide strategic decisions.
 - Complex queries
 - Transactions access a large fraction of the database
 - Data need not be up-to-date
 - E.g., A bank CEO wants to get an approximate total balance of all customers' accounts
 - OLAP transactions need not be serializable with respect to other OLTP transactions



Phenomena caused by Concurrent Transactions

- **dirty read**

- A transaction reads data written by a concurrent uncommitted transaction.

- **nonrepeatable read**

- A transaction runs the same query twice and finds that a data has been modified by another transaction.

- **phantom read**

- A transaction runs the same query twice and finds that the result set has new records due to another recently-committed transaction.

- **serialization anomaly**

- The result of successfully committing a group of transactions is inconsistent with all possible orderings of running those transactions one at a time.

Phantom Phenomenon

- E.g., Transaction 1:
select *ID, name* **from** *instructor* **where** *salary* > 90000
- E.g., Transaction 2:
insert into *instructor* **values** ('11111', 'James', 'Marketing', 100000)
- Suppose
 - T1 starts, finds tuples *salary* > 90000 using index and locks them
 - And then T2 executes.
 - Do T1 and T2 conflict? Does tuple level locking detect the conflict?
 - Instance of the **phantom phenomenon**

Transaction Isolation in SQL-92

- **Serializable** — default
- **Repeatable read** — only committed records to be read.
 - Repeated reads of same record must return same value.
 - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- **Read committed** — only committed records can be read.
 - Successive reads of record may return different (but committed) values.
- **Read uncommitted** — even uncommitted records may be read.

| <i>Isolation Level</i> | Dirty Read | Nonrepeatable Read | Phantom Read | Serialization Anomaly |
|------------------------|------------------------|--------------------|------------------------|-----------------------|
| Read uncommitted | Allowed, but not in PG | Possible | Possible | Possible |
| Read committed | Not possible | Possible | Possible | Possible |
| Repeatable read | Not possible | Not possible | Allowed, but not in PG | Possible |
| Serializable | Not possible | Not possible | Not possible | Not possible |

Transaction Definition in SQL

- In SQL, a transaction begins implicitly.
- A transaction in SQL ends by:
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
 - Implicit commit can be turned off by a database directive
 - E.g., in JDBC -- `connection.setAutoCommit(false);`
- Isolation level can be set at database level
- Isolation level can be changed at start of transaction
 - E.g. In SQL **set transaction isolation level serializable**
 - E.g. in JDBC -- `connection.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE)`



Implementation of Isolation Levels

- Locking
 - Lock on whole database vs lock on items
 - How long to hold lock?
 - Shared vs exclusive locks
- Timestamps
 - Transaction timestamp assigned e.g. when a transaction begins
 - Data items store two timestamps
 - Read timestamp
 - Write timestamp
 - Timestamps are used to detect out of order accesses
- Multiple versions of each data item
 - Allow transactions to read from a “snapshot” of the database