



Multicore Computing

Lecture03 - Pthreads



남 범 석

bnam@skku.edu



POSIX Threads

- This lecture will introduce Pthreads for beginners
 - Will not cover every detail of Pthreads though
- Reference
 - <https://computing.llnl.gov/tutorials/pthreads/>



What is a Thread? (Recap from OS class)

- An independent stream of instructions that can be scheduled to run by the operating system
- This independent flow of control is accomplished because a thread maintains its own:
 - Stack pointer
 - Registers
 - Scheduling properties (such as policy or priority)
 - Set of pending and blocked signals
 - Thread specific data
- Because threads within the same process share resources:
 - Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads.
 - Two pointers having the same value point to the same data.
 - Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer



Overview of Programming Models

- A single-thread process:

```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            dot_product( get_row(a, row), get_col(b, col));
```

- A multi-thread process:

```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            create_thread( dot_product(get_row(a, row),  
                                     get_col(b, col)) );
```





Threads vs. Processes

■ Process

- One address space per process
- Each process has its own data (global variables), stack, heap

■ Thread

- Multiple threads share on address space
 - But its own stack and register context
- Threads within the same address space share data (global variables), heap



What is P(POSIX®)Threads?

- POSIX: Portable Operating System Interface
- Standard threads API supported by most vendors
- Historically, hardware vendors have implemented their own proprietary versions of threads
 - Substantially different from each other
 - Make it difficult for programmers to develop portable threaded apps
- The need for a standardized programming interface
 - For UNIX systems, this interface has been specified by the IEEE POSIX 1003.1c standard (1995)
 - POSIX: Portable Operating System Interface
 - Implementations adhering to this standard are referred to as POSIX threads, or Pthreads
- Pthreads are defined as a set of C language programming types and procedure calls
 - Implemented with a pthread.h header/include file and a thread library



Why Pthreads?

- Mainly for higher performance gains when exploiting parallelism
 - Much more efficient than creating and managing processes
- Example: creating 50,000 processes/threads (units in seconds)

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6



The Pthreads API

- The subroutines of the Pthreads API can be informally grouped into the following major groups:
 - Thread management:
 - Routines that work directly on threads - creating, detaching, joining, etc
 - Mutexes:
 - Routines that deal with synchronization, called a "mutex", which is an abbreviation for "mutual exclusion"
 - Condition variables:
 - Routines that address communications between threads that share a mutex.
 - Synchronization:
 - Routines that manage read/write locks and barriers.



Pthreads: Naming Conventions

- Naming conventions:
 - All identifiers in the threads library begin with pthread_.

Routine Prefix	Functional Group
pthread_	Threads themselves and miscellaneous subroutines
pthread_attr_	Thread attributes objects
pthread_mutex_	Mutexes
pthread_mutexattr_	Mutex attributes objects.
pthread_cond_	Condition variables
pthread_condattr_	Condition attributes objects
pthread_key_	Thread-specific data keys
pthread_rwlock_	Read/write locks
pthread_barrier_	Synchronization barriers



Thread Creation and Destruction

- Typical Pthreads code would look like below

```
#include <pthread.h>
#define NT 5
int err;

void main()
{
    pthread_t tid[NT]; // array of thread IDs, one per thread
    for (i=0; i<NT; i++)
        err=pthread_create(&tid[i], NULL, printHello, (void*) i);

    for (i=0; i<NT; i++)
        err=pthread_join(tid[i], (void **)&status[i]);
}
```

- Threads are created in the first loop
 - Parent → the creating thread, children → created threads
- The parent thread waits for its children to complete



Creating a Pthread

```
int pthread_create(  
    pthread_t* thread /* out */ ,  
    const pthread_attr_t* attr /* in */ ,  
    void* (*start_routine) (void*) /* in */ ,  
    void* arg /* in */ ) ;
```

Specify thread handle (allocated before calling)

Specify the attributes of a creating thread

The function that the thread is to run

Pointer to the argument passed to the function ***start_routine***



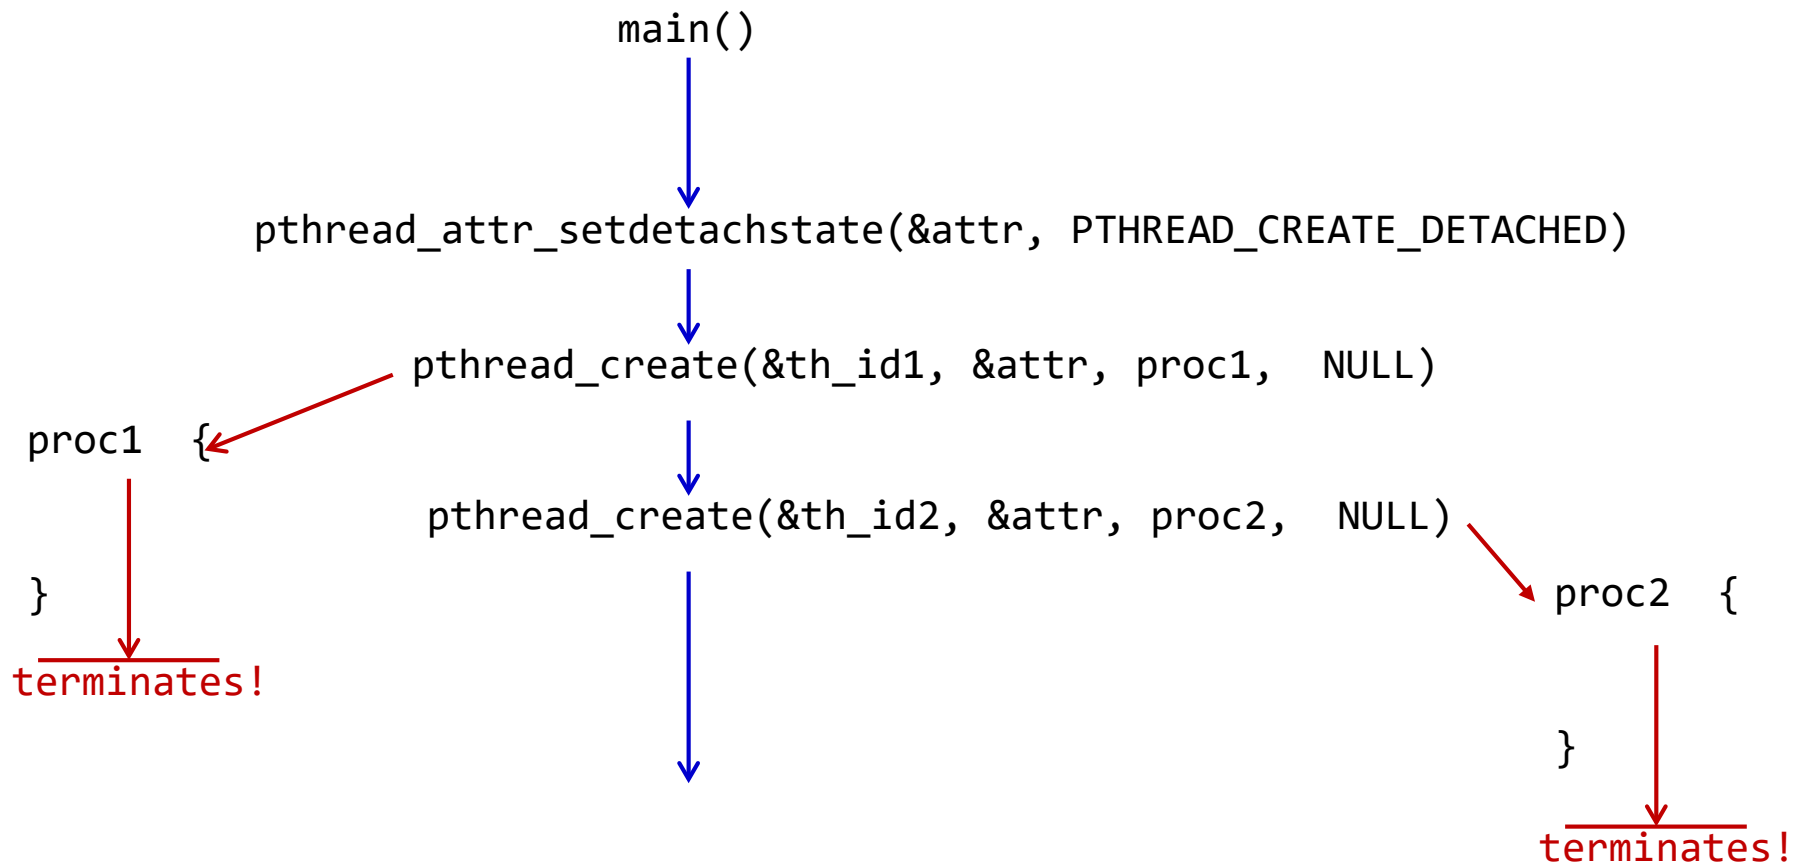
Pthread Attributes

- Stack size
- Detach state
 - PTHREAD_CREATE_DETACHED, PTHREAD_CREATE_JOINABLE
 - Release resources at termination (detached) or retain (joinable)
- Scheduling policy
 - SCHED_OTHER: standard policy
 - SCHED_FIFO, SCHED_RR
- Scheduling parameters
 - Priority only
- Inherit scheduling policy
 - PTHREAD_INHERIT_SCHED, PTHREAD_EXPLICIT_SCHED
- Thread scheduling scope
 - PTHREAD_SCOPE_SYSTEM, PTHREAD_SCOPE_PROCESS
- Special functions exist for getting/setting each attribute
 - `int pthread_attr_setstack_size(pthread_attr_t* attr, size_t stacksize)`



Pthreads – detached thread

- `pthread_attr_setdetachstate`
 - Detached threads are not joinable

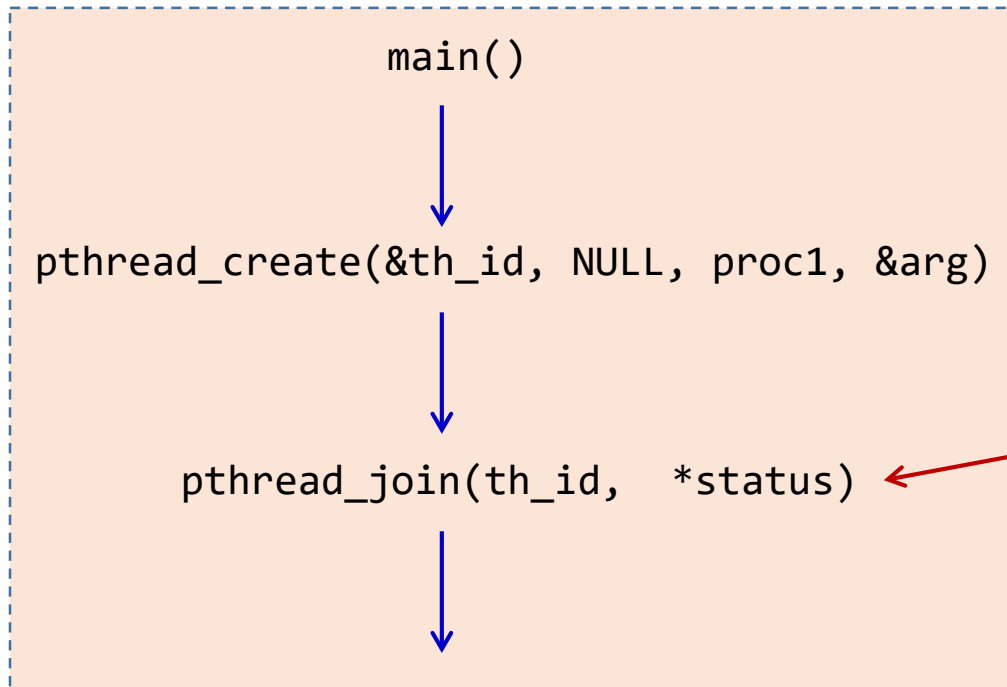


Waiting a Pthread

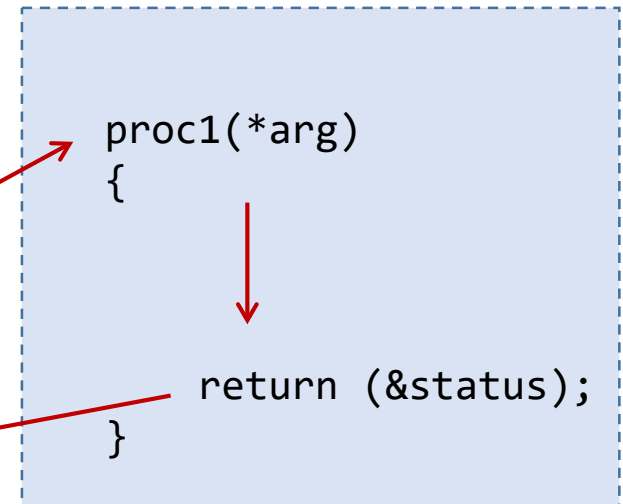
```
int pthread_join (
    pthread_t thread, void **ptr);
```

- Returns after a specified thread terminates
- ptr stores return code of a terminating thread

thread 1



thread 2



Exiting a Pthread

- 4 ways to exit threads
 - Thread will naturally exit after starting thread function returns
 - Thread itself can exit by calling `pthread_exit()`
 - Other threads can terminate a thread by calling `pthread_cancel()`
 - A specified thread terminates when it reaches a **cancellation point**
 - Thread exits if the process that owns the thread exits
- APIs
 - `void pthread_exit (void *retval);`
 - `int pthread_cancel (pthread_t thread)`



Pthreads – compilation

- Pthreads are supported by almost all compilers
 - GNU Compiler
 - `gcc -Wall -o hello hello.c -lpthread`
 - `-lxxx` : specifies which static library to link
 - `-Wall` : specifies to print out all types of warnings

Thread Argument Passing

```
struct thread_data{
    int  thread_id;
    int  sum;
    char *message;
};

struct thread_data thread_data_array[NUM_THREADS];

void *PrintHello(void *threadarg)
{
    struct thread_data *my_data;
    ...
    my_data = (struct thread_data *) threadarg;
    taskid = my_data->thread_id;
    sum = my_data->sum;
    hello_msg = my_data->message;
    ...
}

int main (int argc, char *argv[])
{
    ...
    thread_data_array[t].thread_id = t;
    thread_data_array[t].sum = sum;
    thread_data_array[t].message = messages[t];
    rc = pthread_create(&threads[t], NULL, PrintHello,
        (void *) &thread_data_array[t]);
    ...
}
```

- We often want to pass multiple parameters to threads when creating them
 - create a struct which contains all of the arguments, and pass a pointer to that struct
- Each thread needs to have different parameter values
 - Allocate an array of the struct whose length equals to the thread count
 - Each thread references its own parameters using its thread ID as an index to the array



Synchronization

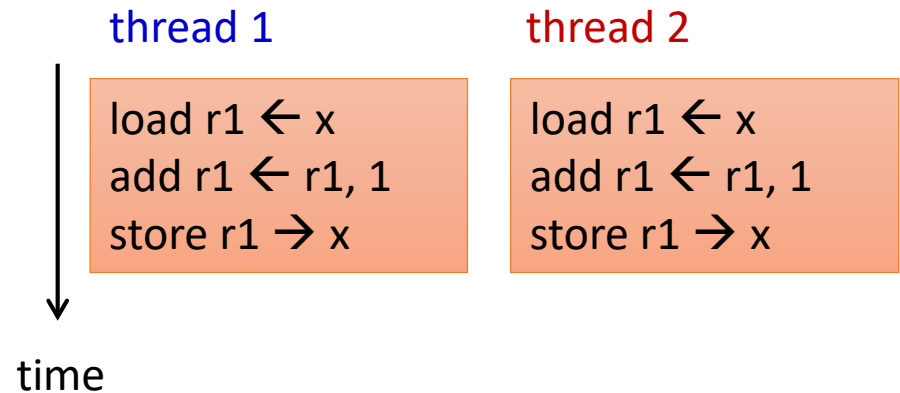
- Accessing shared data
 - Example: two threads increase the same variable x

```
int x = 0;
void* inc () {
    x = x + 1;
    return NULL;
}

main()
{
    pthread_create(&th1, NULL, inc, NULL);
    pthread_create(&th2, NULL, inc, NULL);

    pthread_join(th1, NULL);
    pthread_join(th2, NULL);

    printf("x = %d\n", x);
}
```



Critical Sections

- Critical section
 - Need to guarantee that one process (thread) can access a certain resource at a time
 - Implemented mechanism is known as "mutual exclusion"
- Locks
 - Simple mechanism for mutual exclusion
 - A lock can have only two values
 - 1 – a thread entered the critical section
 - 0 – no thread is in the critical section
 - Acquire the lock before entering the critical section (set to 1)
 - Release the lock after leaving the critical section (set to 0)

Acquiring Locks

- Does the following simple C code work?

```
void lock(lock_var) {  
    while (lock_var != 0);  
    lock_var = 1;  
}  
  
void unlock(lock_var) {  
    lock_var = 0;  
}
```

- Special *atomic* instruction should be used
 - Pthread library provides APIs

Mutual Exclusion (Mutex)

- Pthreads provides mutex to avoid race conditions

```
pthread_mutex_t lock=PTHREAD_MUTEX_INITIALIZER;
```

```
...
```

```
pthread_mutex_lock(&lock);
```

```
// critical section
```

```
pthread_mutex_unlock(&lock);
```

- pthread_mutex_lock
 - A thread will wait until it can acquire the lock
- pthread_mutex_unlock
 - If multiple threads are waiting, only one thread is selected to receive the lock
 - Only the thread that acquires the lock can unlock it



Example: Sum an array using Pthreads

```
#include <pthread.h>

int a[array_size];
int global_index = 0;
int sum = 0;
pthread_mutex_t mutex1;

void *do_work(void *tid)
{
    int i, start, mytid, end;
    int local_sum=0;
    mytid = (int) tid;
    start = mytid*array_size/no_threads;
    end = start+array_size/no_threads;
    for (i=start; i<end; i++) {
        local_sum = a[i];
    }

    pthread_mutex_lock(&mutex1);
    sum += local_sum;
    pthread_mutex_unlock(&mutex1);

    pthread_exit(NULL);
}
```

```
void main() {
    int i;
    pthread_t thread[no_threads];
    pthread_attr_t attr;

    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr,
                                PTHREAD_CREATE_JOINABLE);

    ...

    for (i = 0; i < no_threads; i++)
        pthread_create(&thread[i], &attr,
                      do_work, (void*) i);

    for (i = 0; i < no_threads; i++)
        pthread_join(thread[i], NULL);

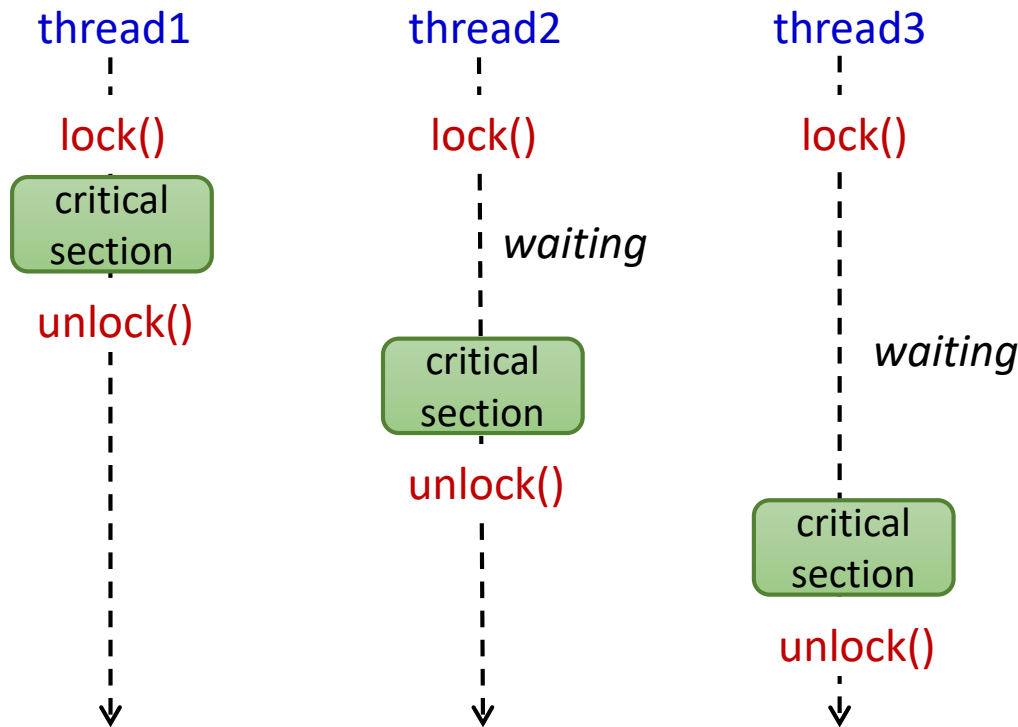
    printf("The sum of 1 to %i is %d\n",
           array_size, sum);

    pthread_attr_destroy(&attr);
    pthread_exit(NULL);
}
```



Serialization

- Critical sections serialize the code execution
 - Too many or large critical sections can slow down the performance – sequential code may run faster



Alleviating Locking Overhead

```
int pthread_mutex_trylock (  
    pthread_mutex_t *mutex_lock);
```

- Reduce overhead by overlapping computation with waiting
 - Acquires lock if unlocked
 - Returns **EBUSY** if locked



Condition Variables

- Wait until a condition is satisfied
 - A global variable is used to indicate condition (predicate value)
- Three variables are linked all together
 - mutex lock, condition variable, predicate

thread1

```
action() {  
    ...  
    mutex_lock(&lock);  
    while (predicate == 0) // test predicate  
        cond_wait(&cond, &lock);  
    mutex_unlock(&lock);  
    // perform action  
    ...  
}
```

thread2

```
signal() {  
    ...  
    mutex_lock(&lock);  
    predicate = 1; // set predicate  
    cond_signal(&cond);  
    mutex_unlock(&lock);  
    ...  
}
```

- When a thread waits using `cond_wait`, associated mutex is unlocked
- If a thread is signaled, returns after acquiring the mutex lock



Pthread Condition Variable API

- Initialize and destroy

```
int pthread_cond_init(pthread_cond_t *cond,  
                      const pthread_condattr_t *attr);  
int pthread_cond_destroy(pthread_cond_t *cond);
```

- Wait for a condition

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_timedwait(pthread_cond_t *cond,  
                            pthread_mutex_t *mutex,  
                            const struct timespec *wtime);
```


- Signal one or all waiting threads

```
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);
```



Bounded Buffer Example

- One or more threads put items into a circular buffer
 - They need to wait if the buffer is full
- Other threads remove items from the same buffer
 - They need to wait if the buffer is empty



Bounded Buffer Example

```
1  pthread_mutex_t lock=PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t nonempty=PTHREAD_COND_INITIALIZER;
3  pthread_cond_t nonfull=PTHREAD_COND_INITIALIZER;
4  Item buffer[SIZE];
5  int put=0;                      // Buff index for next insert
6  int get=0;                      // Buff index for next remove
7
8  void insert(Item x)              // Producer thread
9  {
10     pthread_mutex_lock(&lock);
11     while((put>get&&(put-get)==SIZE-1)|| // While buffer is
12           (put<get&&(put+get)==SIZE-1)) // full
13     {
14         pthread_cond_wait(&nonfull, &lock);
15     }
16     buffer[put]=x;
17     put=(put+1)%SIZE;
18     pthread_cond_signal(&nonempty);
19     pthread_mutex_unlock(&lock);
20 }
21
22 Item remove()                    // Consumer thread
23 {
24     Item x;
25     pthread_mutex_lock(&lock);
26     while(put==get)               // While buffer is empty
27     {
28         pthread_cond_wait(&nonempty, &lock);
29     }
30     x=buffer[get];
31     get=(get+1)%SIZE;
32     pthread_cond_signal(&nonfull);
33     pthread_mutex_unlock(&lock);
34     return x;
35 }
```

- Two condition variables are used
 - nonempty
 - nonfull
- Two cursors
 - put: next empty location
 - get: points to next element to remove
- A mutex variable to protect access to CVs by multiple threads



Protecting Condition Variables

```
1 pthread_mutex_t lock=PTHREAD_MUTEX_INITIALIZER;
2 pthread_cond_t nonempty=PTHREAD_COND_INITIALIZER;
3 pthread_cond_t nonfull=PTHREAD_COND_INITIALIZER;
4 Item buffer[SIZE];
5 int put=0; // Buff index for next insert
6 int get=0; // Buff index for next remove
7
8 void insert(Item x) // Producer thread
9 {
10  pthread_mutex_lock(&lock);
11  while((put>get&&(put-get)==SIZE-1)|| // While buffer is
12        (put<get&&(put+get)==SIZE-1)) // full
13  {
14    pthread_cond_wait(&nonfull, &lock);
15  }
16  buffer[put]=x;
17  put=(put+1)%SIZE;
18  pthread_cond_signal(&nonempty);
19  pthread_mutex_unlock(&lock);
20 }
21
22 Item remove() // Consumer thread
23 {
24  Item x;
25  pthread_mutex_lock(&lock);
26  while(put==get) // While buffer is empty
27  {
28    pthread_cond_wait(&nonempty, &lock);
29  }
30  x=buffer[get];
31  get=(get+1)%SIZE;
32  pthread_cond_signal(&nonfull);
33  pthread_mutex_unlock(&lock);
34  return x;
35 }
```

Why do we need the while loop in Line 11?

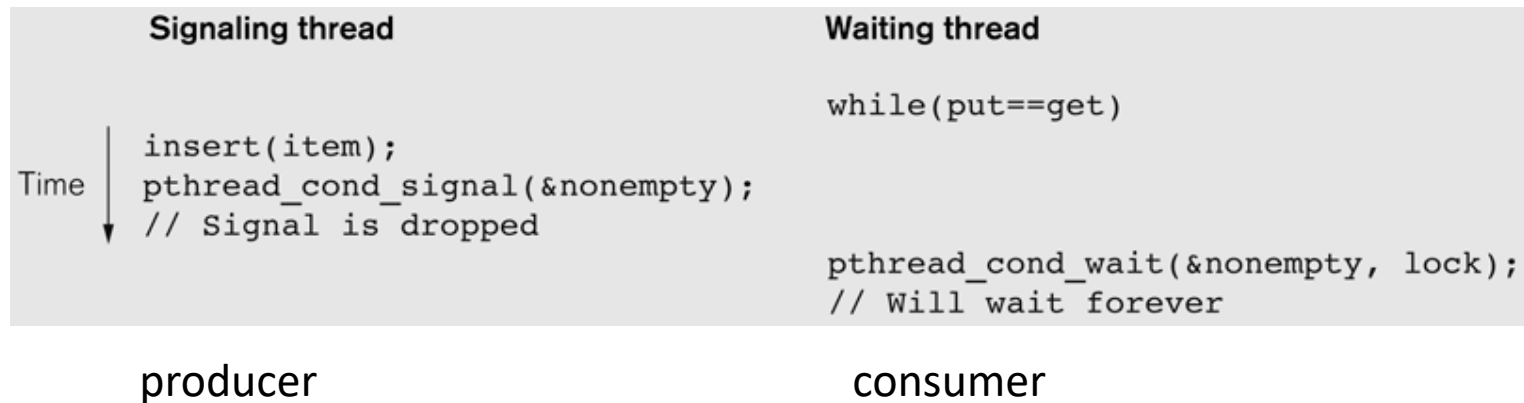
- At the time of the signal, the buffer is not full, but when any particular thread acquires the mutex, the buffer may have become full again
- In this case, the thread should call `pthread_cond_wait()` again

Why do we need to protect `pthread_cond_signal()` in line 18 and 32?



Protecting Condition Variables

- Calls to `pthread_cond_signal()` should be protected
 - To avoid dropped signals
 - Consider a consumer acquires the lock and finds the buffer is empty, so it executes `pthread_cond_wait()`.
 - If the producer does not protect the call to `pthread_cond_signal()` with a mutex, it could insert an item into the buffer immediately after the waiting thread found it empty.
 - If the producer then signals the buffer is nonempty before the waiting thread executes the call to `pthread_cond_wait()`, the signal will be dropped and the consumer thread will not realize that the buffer is actually not empty.



Thread-Specific Data

- Goal: associate some data with a thread
- Choices
 - Pass data as arguments to functions
 - Store data in a shared array indexed by thread id
 - Use thread-specific data API

- API

```
int pthread_key_create(pthread_key_t *key,  
    void (*destr_function) (void *));  
int pthread_key_delete(pthread_key_t key);  
int pthread_setspecific(pthread_key_t key,  
    const void *pointer);  
void * pthread_getspecific(pthread_key_t key);
```



Example: Thread-Specific Data

```
/* Key for the thread-specific buffer */
static pthread_key_t buffer_key;

/* Free each thread-specific buffer when associated
thread terminates */
static void buffer_destroy(void * buf){
    free(buf);
}

int main(){
    /* Initialize key, only once */
    pthread_key_create(&buffer_key, buffer_destroy);
    ...
}
```



Example: Thread-Specific Data

```
void* thread_func(){
    /* get the thread-specific buffer */
    if( (buf = pthread_getspecific(buffer_key)) == NULL){
        /* Allocate a thread-specific buffer */
        buf = malloc(100);
        pthread_setspecific(buffer_key, buf);
    }
    memcpy(buf, ...) // do some work
}

int main(){
    ...
    /* delete key, only once */
    pthread_key_delete(&buffer_key);
}
```



Other Useful APIs in Pthread

- Read-write lock
- Spinlock
- Barrier
- Caveat
 - Not all Pthread libraries provide those
 - One can implement those functions using mutexes, condition variables, and other variables



Read-Write Locks

- Useful for applications having a frequently read but infrequently written data structure
- Provide a critical section that
 - Multiple reader can be in the critical section simultaneously
 - Increasing concurrency of execution
 - One writer can be in the critical section at a time
 - Avoiding race condition



Pthread Read-Write Mutex API

- Initialize and destroy a rwlock
 - `pthread_rwlock_t` rwlock = PTHREAD_RWLOCK_INITIALIZER
 - `pthread_rwlock_init`(pthread_rwlock_t *rwlock, pthread_rwlockattr_t *attr)
 - `pthread_rwlock_destroy`(pthread_rwlock_t* rwlock)
- Read locking
 - `pthread_rwlock_rdlock`(pthread_rwlock_t* rwlock)
- Write locking
 - `pthread_rwlock_wrlock`(pthread_rwlock_t* rwlock)
- Unlocking
 - `pthread_rwlock_unlock`(pthread_rwlock_t* rwlock)



Pthread Spinlock API

- Pthread mutex is blocking lock
 - Inefficient when a critical section is short
 - Due to management cost for blocking/waking-up a thread
- Busy-waiting lock (e.g., spinlock)
 - More efficient when a critical section is short
 - Spinning a few cycles is faster than blocking and waking up a thread
- Pthread library provides spinlock-based mutual exclusion

`pthread_spinlock_t`

`pthread_spin_init(pthread_spinlock_t* spinlock, int nr_shared)`

`pthread_spin_lock(pthread_spinlock_t* spinlock)`


`pthread_spin_unlock(pthread_spinlock_t* spinlock)`



Pthread Barrier API

■ Barrier

- An execution synchronization point of threads
- Wait until all thread reach the point



■ API

- `pthread_barrier_init(pthread_barrier_t* barrier, pthread_barrierattr_t* attr, int value)`
 - Initialize a barrier
 - The integer value specifies the number of threads to synchronize
- `pthread_barrier_wait(pthread_barrier_t* barrier)`
 - Waits until the specified number of threads arrives at the barrier

