



Multicore Computing

Lecture04 - OpenMP



남 범 석
bnam@skku.edu





Topic Overview

- Introduction to OpenMP
- OpenMP directives
 - Concurrency control
 - parallel, for, sections
 - Synchronization
 - reduction, barrier, single, master, critical, atomic, ordered, ...
 - Data handling
 - private, shared, firstprivate, lastprivate, threadprivate, ...
- OpenMP library APIs
- Environment variables



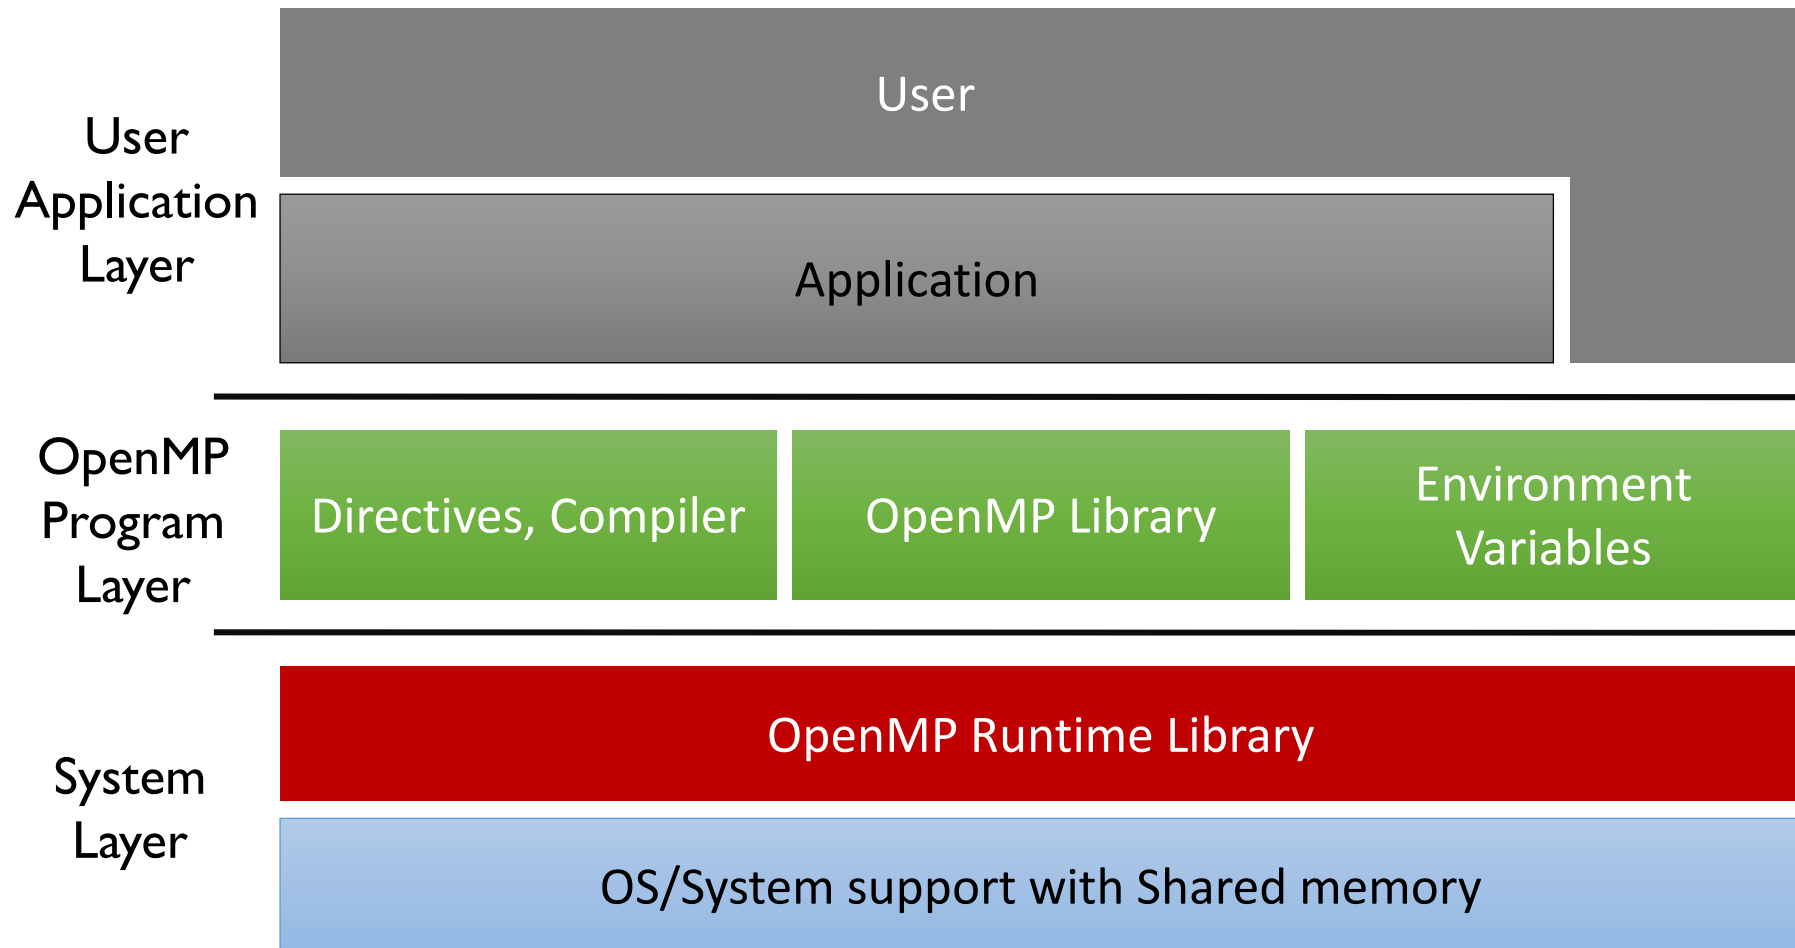


OpenMP

- Open specifications for Multi Processing
- A standard for directive-based Parallel Programming
 - Shared-address space programming
 - FORTRAN, C, and C++
 - Support concurrency, synchronization, and data handling
 - Obviate the need for explicitly setting up mutexes, condition variables, data scope, and initialization



OpenMP Solution Stack





Parallel Programming Practice

■ Current

- Start with a parallel algorithm
- Implement, keeping in mind
 - Data races
 - Synchronization
 - Threading syntax
- Test & Debug
- Debug

■ Ideal way

- Start with some algorithm
- Implement serially, ignoring
 - Data races
 - Synchronization
 - Threading syntax
- Test & Debug
- Auto-magically parallelize



Implementation on Shared Memory

- Thread Library
 - Library calls
 - Low level programming
 - Explicit thread creation & work assignment
 - Explicit handling of synchronization
 - Parallelism expression
 - Task: create/join thread
 - Data: detailed programming
 - Design concurrent version from the start
- OpenMP
 - **Compiler directives**
 - Higher abstraction
 - Compilers convert code to use OpenMP library, which is actually implemented with thread APIs
 - Parallelism expression
 - Task: task/taskwait, parallel sections
 - Data: parallel for
 - Incremental development
 - Start with sequential version
 - Insert necessary directives



OpenMP Example

- Pragmas (compiler directives)

Parallel section
executed by
all threads

```
1  int count3s()
2  {
3      int i, count_p;
4      count=0;
5      #pragma omp parallel shared(array, count, length)\
6          private(count_p)
7      {
8          count_p=0;
9          #pragma omp parallel for private(i)
10         for(i=0; i<length; i++)
11         {
12             if(array[i]==3)
13             {
14                 count_p++;
15             }
16         }
17         #pragma omp critical
18         {
19             count+=count_p;
20         }
21     }
22     return count;
23 }
```

Fork a set of threads

Join threads and resume serial code



OpenMP Example

- Threaded functions
 - Exploit data parallelism

```
node A[N], B[N];

main() {
    for (i=0; i<nproc; i++)
        thread_create(par_distance);
    for (i=0; i<nproc; i++)
        thread_join();
}

void par_distance() {
    tid = thread_id();  n = ceiling(N/nproc);
    s = tid * n;        e = MIN((tid+1)*n, N);
    for (i=s; i<e; i++)
        for (j=0; j<N; j++)
            C[i][j] = distance(A[i], B[j]);
}
```

- Parallel loops
 - Exploit data parallelism

```
node A[N], B[N];

#pragma omp parallel for
for (i=0; i<N; i++)
    for (j=0; j<N; j++)
        C[i][j] = distance(A[i], B[j]);
```





Compiler Directives

- Appear as comments in your source code
 - Ignored by compilers unless you tell them otherwise
- OpenMP compiler directives are used for various purposes:
 - Spawning a parallel region
 - Dividing blocks of code among threads
 - Distributing loop iterations between threads
 - Serializing sections of code
 - Synchronization of work among threads
- Syntax:
`#pragma omp <specifications>`





Runtime Library Routines

- Routines are used for a variety of purposes
 - Setting and querying the number of threads
 - Querying a thread's unique identifier (thread ID), a thread's ancestor's identifier, the thread team size
 - Setting and querying the dynamic threads feature
 - Querying if in a parallel region, and at what level
 - Setting and querying nested parallelism
 - Setting, initializing and terminating locks and nested locks
 - Querying wall clock time and resolution
- Example

```
int omp_get_num_threads(void)
```





Environment Variables


- Environment variables can be used to control such things as:
 - Setting the number of threads
 - Specifying how loop iterations are divided
 - Binding threads to processors
 - Enabling/disabling nested parallelism; setting the maximum levels of nested parallelism
 - Enabling/disabling dynamic threads
 - Setting thread stack size
 - Setting thread wait policy

- Example
 - `export OMP_NUM_THREADS=8`



OpenMP Programming Model

- Fork-join model
 - Thread pool
 - Implicit barrier
 - `#pragma omp`
 - `parallel for`
 - `parallel sections`



- Data scoping semantics are somewhat complicated
 - `private`, `shared`, `copyin`, `firstprivate`, `lastprivate`, `copyprivate`, `threadprivate`, ...
 - Implicit rules,...

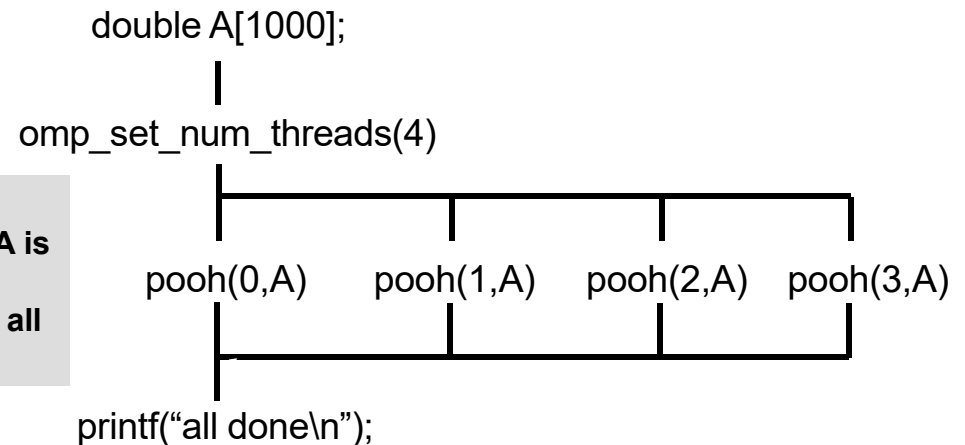


How to set number of threads?

- Setting of the NUM_THREADS clause
- Or, use of the `omp_set_num_threads()` library function
- Or, setting of the `OMP_NUM_THREADS` environment variable

```
double A[1000];  
  
#pragma omp parallel num_threads(4)  
{  
    int ID = omp_get_thread_num();  
    Pooh(ID,A);  
}  
  
printf("all done\n");
```

**A single
copy of A is
shared
between all
threads.**



Hello World in OpenMP

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
} /* Hello */
```



Hello World in OpenMP – con't

- Compile
 - `#gcc -g -Wall -fopenmp -o omp_hello omp_hello.c`
- Run
 - `#./omp_hello 4`

Possible outcomes

Hello from thread 0 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 3 of 4

Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4
Hello from thread 3 of 4

Hello from thread 3 of 4
Hello from thread 1 of 4
Hello from thread 2 of 4
Hello from thread 0 of 4



Pragmas

- Special compiler directives
 - **#pragma**
 - Provides extension to the basic C (or C++)
 - Compilers that don't support the pragmas **ignore** them
- OpenMP pragmas
 - #pragma omp directive [clause list]**
/* structured block */
 - Directives specify actions OpenMP supports
 - Additional clauses follow the directive
 - Parallel directive
 - #pragma omp parallel [clause list]**
 - Most basic parallel directive in OpenMP



parallel Directive

`#pragma omp parallel [clause list]`

- Possible clauses
 - Conditional Parallelization
 - `if` (scalar expression)
 - Determines whether to create threads or not
 - Degree of Concurrency
 - `num_threads` (integer expression)
 - Specifies the number of threads that are created.
 - Data Handling
 - `private` (variable list)
 - Variables local to each thread
 - `firstprivate` (variable list)
 - Variables are initialized to corresponding values before the directive
 - `lastprivate`: (variable list)
 - PRIVATE + copy from the last thread execution
 - `shared` (variable list)
 - Variables are shared across all the threads.
 - `default` (shared|private|none)
 - Default data handling specifier



parallel Directive

```
1  int count3s()
2  {
3      int i, count_p;
4      count=0;
5      #pragma omp parallel shared(array, count, length)\
6          private(count_p)
7      {
8          count_p=0;
9          #pragma omp parallel for private(i)
10         for(i=0; i<length; i++)
11         {
12             if(array[i]==3)
13             {
14                 count_p++;
15             }
16         }
17         #pragma omp critical
18         {
19             count+=count_p;
20         }
21     }
22     return count;
23 }
```



parallel Directive

`#pragma omp parallel [clause list]`

- Possible clauses
 - Conditional Parallelization
 - `if` (scalar expression)
 - Determines whether to create threads or not
 - Degree of Concurrency
 - `num_threads` (integer expression)
 - Specifies the number of threads that are created.
 - Data Handling
 - `private` (variable list)
 - Variables local to each thread
 - `firstprivate` (variable list)
 - Variables are initialized to corresponding values before the directive
 - `lastprivate`: (variable list)
 - PRIVATE + copy from the last thread execution
 - `shared` (variable list)
 - Variables are shared across all the threads.
 - `default` (shared|private|none)
 - Default data handling specifier



Example of **parallel** Directive

```
#pragma omp parallel if (is_parallel== 1) num_threads(8) \  
    private (a) shared (b) firstprivate(c) default(none) {  
    /* structured block */  
}
```

- **if** (is_parallel==1) **num_threads** (8)
 - If the value of the variable `is_parallel` equals one, eight threads are created.
- **private** (a)
 - Threads get private copy of variable a
- **firstprivate** (c)
 - private copy + initialization
 - The value of each copy of `c` is initialized to the value of `c` before the parallel directive.
- **shared** (b)
 - Threads share a single copy of variable b.
- **default** (none)
 - Default scope of variables are none
 - Compile error when not all variables are specified as **shared** or **private**



for Directive

- Split parallel iteration spaces (i.e., loop) across threads


- Implicit barrier at the end of a loop

- General form

- `#pragma omp for [clause list]`
 - `/* for loop */`

- Possible clauses

- `private`, `firstprivate`, `lastprivate`, `reduction`, `schedule`, `nowait`, and `ordered`.



for Directive example

- OpenMP shortcut: Put the “parallel” and the worksharing directive on the same line

```
double res[MAX]; int I;  
#pragma omp parallel{  
    #pragma omp for  
    for (i=0;i<MAX;i++){  
        res[i] = huge();  
    }  
}
```

```
double res[MAX]; int I;  
#pragma omp parallel for  
    for (i=0;i<MAX;i++){  
        res[i] = huge();  
    }  
}
```

↙ ↘
These are equivalent



Data Sharing: Private clause

- `private(var)` creates a new local copy of `var` for each thread
 - The value is uninitialized

```
void wrong() {  
    int tmp = 0;  
    #pragma omp parallel for private(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```

Q: What is wrong in this code?



Data Sharing: FirstPrivate clause

- firstprivate(var) is a special case of private
 - Initializes each private copy with the corresponding value from the master thread

```
void useless() {  
    int tmp = 0;  
    #pragma omp parallel for firstprivate(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```

Each thread gets its own tmp with an initial value of 0, but

Q: What is wrong in this code?



Data Sharing: LastPrivate clause

- Lastprivate passes the value of a private from the last iteration to a global variable.

```
void closer() {  
    int tmp = 0;  
    #pragma omp parallel for firstprivate(tmp) lastprivate(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```



Reduction

- How do we handle this case?

```
double ave=0.0, A[MAX];  
int i;  
for (i=0;i< MAX; i++) {  
    ave + = A[i];  
}  
ave = ave/MAX;
```

- We are combining values into a single accumulation variable (ave) ... there is a true dependence between loop iterations that can't be trivially removed
- This is a very common situation ... it is called a "reduction".
- Support for reduction operations is included in most parallel programming environments.



reduction Clause

- A **reduction**

- `#pragma omp parallel for reduction(+: sum) num_threads(8) {
 for (row = 0; row < Rows; row++) sum += val;
}`

- Applies the same reduction operator to a sequence of operands to get a single result
 - Reduction operators → **+, *, -, &, |, ^, &&, ||**
 - Commutative and associative operators can provide correct results
 - All of the intermediate results of the operation should be stored in the same variable: the reduction variable

- Reduction clause in OpenMP

- **reduction**(**<operator>**: **<variable list>**)
 - The variables in the list are implicitly specified as being private to threads.



Example: estimating Pi

- Estimating Pi using Monte Carlo method




Image credit: wikipedia.org



Estimating Pi using OpenMP

```
/* *****  
An OpenMP version of a threaded program to compute PI.  
***** */  
  
#pragma omp parallel default(private) shared(npoints) \  
    reduction(+: sum) num_threads(8)  
{  
    num_threads = omp_get_num_threads();  
    sample_points_per_thread = npoints / num_threads;  
    sum = 0;  
    for (i = 0; i < sample_points_per_thread; i++) {  
        rand_no_x = (double)rand_r(&seed)/RAND_MAX;  
        rand_no_y = (double)rand_r(&seed)/RAND_MAX;  
        if ( x * x + y * y < 1.0)  
            sum ++;  
    }  
}
```



Estimating Pi using Pthreads (1)

```
#include <pthread.h>
#include <stdlib.h>
#define MAX_THREADS 512
void *compute_pi (void *);
....
main() {
    ...
    pthread_t p_threads[MAX_THREADS];
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    for (i=0; i< num_threads; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void *) &hits[i]);
    }
    for (i=0; i< num_threads; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```



Estimating Pi using Pthreads (2)

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double x, y;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        x =(double)rand_r(&seed)/RAND_MAX;
        y =(double)rand_r(&seed)/RAND_MAX;
        if ( x * x + y * y < 1.0 )
            local_hits ++;
        seed *= i;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```



Example: Estimating Pi using **for**

```
#pragma omp parallel default(private) \
    shared(npoints) reduction(+: sum) num_threads(8)
{
    sum = 0;
    #pragma omp for
    for (i = 0; i < sample_points_per_thread; i++) {
        rand_no_x = (double)rand_r(&seed) / RAND_MAX;
        rand_no_y = (double)rand_r(&seed) / RAND_MAX;
        if ( x * x + y * y < 1.0)
            sum ++;
    }
}
```



