



Multicore Computing

Lecture05 - OpenMP Part II




남 범 석

bnam@skku.edu



Assigning Iterations to Threads

- `parallel for` directive
 - Basic partitioning policy → block partitioning



- Is this optimal?
 - **Yes**, when each iteration takes equal time
 - **No**, when each iteration takes different time
 - ex) larger index takes longer time




Assigning Iterations to Threads


- Example of $f(i)$
 - $f(i)$ calls \sin function i times
 - $\text{Time}(f(i))$ is linear to i

```
float A[100][100];  
for(int i = 0; i < 100; i++) {  
    for(int j = 0; j < i; j++) {  
        A[i][j] = 1.0f;  
    }  
}
```

- Block partitioning



- Cyclic partitioning




Assigning Iterations to Threads

- Example of $f(i)$
 - $f(i)$ calls sin function i times
 - Time($f(i)$) is linear to i
 - $n = 10,000$

```
float A[100][100];  
for(int i = 0; i < 100; i++) {  
    for(int j = 0; j < i; j++) {  
        A[i][j] = 1.0f;  
    }  
}
```

	One thread	Two threads	
		Block partitioning	Cyclic partitioning
Run-time	3.67s	2.77s	1.84s
Speed-up	1x	1.33x	1.99x

- Scheduling (load balancing) is important



Schedule Clause

■ Format

- `#pragma omp parallel for schedule(type [, chunk])`
 - *type* := static, dynamic, guided, runtime
 - *chunk* := positive integer (# of iterations)
- **static**
 - Divide iterations by *chunk* (near equal in size by default)
 - Statically assign threads in a round-robin fashion
- **dynamic**
 - Divide iterations by *chunk* (1 by default)
 - Dynamically assign a chunk to an idle thread (master/worker)
- **guided**
 - Chunk size is reduced in an exponentially decreasing manner
 - Dynamically assign a chunk to an idle thread (master/worker)
 - Minimum chunk size is specified by *chunk* (1 by default)
- **runtime**
 - Determined at runtime with OMP_SCHEDULE environment variable
 - E.g., export OMP_SCHEDULE="dynamic,5"



Schedule Clause (Illustration)

- Dividing iteration space
 - **Static** schedule on iteration space



- **Dynamic** schedule on iteration space (master/worker)



- **Guided** schedule on iteration space (master/worker)



The Static Schedule Type

#pragma omp parallel for `schedule(static, chunksize)`

- Static schedule on iteration space



- Example

- 12 iterations(0, 1, ..., 11) and 3 threads

`schedule(static, 1)`

Thread 0 : 0, 3, 6, 9

Thread 1 : 1, 4, 7, 10

Thread 2 : 2, 5, 8, 11

`schedule(static, 2)`

Thread 0 : 0, 1, 6, 7

Thread 1 : 2, 3, 8, 9

Thread 2 : 4, 5, 10, 11

`schedule(static, 4)`

Thread 0 : 0, 1, 2, 3

Thread 1 : 4, 5, 6, 7

Thread 2 : 8, 9, 10, 11



The Dynamic Schedule Type

```
#pragma omp parallel for schedule(dynamic, chunksize)
```

- **Dynamic** schedule on iteration space (master/worker)



- The iterations are broken up into chunks of **chunksize** consecutive iterations
- Each thread executes a chunk
- When a thread finishes a chunk, it requests another one from the run-time system
- This continues until all the iterations are completed
- The **chunksize** is 1 by default



The Guided Schedule Type

```
#pragma omp parallel for schedule(guided, chunksize)
```

- **Guided** schedule on iteration space (master/worker)




- Initial size of chunk is smaller than $\# \text{ of iterations} / \# \text{ of threads}$
- Each thread executes a chunk, and when a thread finishes a chunk, it requests another one from runtime system
- The **size of the new chunk decreases exponentially**
- If no **chunksize** is specified, the size of the chunks decreases down to 1
- If **chunksize** is specified, it decreases down to **chunksize**
 - The very last chunk can be smaller than **chunksize**




nowait Clause

- Worksharing for loops end with an implicit barrier
- Often, less synchronization is appropriate
 - When a series of for directives are in a `parallel` construct
- `nowait` clause
 - Used with a `for` directive
 - Avoids implicit barrier at the end of `for`

```
#pragma omp for  
for ( ... )  
#pragma omp for  
for ( ... )
```



```
#pragma omp for nowait  
for ( ... )  
#pragma omp for  
for ( ... )
```



sections Directive

- Non-iterative parallel task assignment using the sections directive.

- General form

```
#pragma omp sections [clause list]
{
    #pragma omp section
    /* structured block */

    #pragma omp section
    /* structured block */

    ...
}
```


- Possible clauses
 - nowait, shared, private, ...
- Implicit barrier at the end of sections




sections Directive: Example

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            taskA();
        }
        #pragma omp section
        {
            taskB();
        }
        #pragma omp section
        {
            taskC();
        }
    }
}
```

n threads available




1 thread available



Merging Directives (1)

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        {
            taskA();
        }
        #pragma omp section
        {
            taskB();
        }
        #pragma omp section
        {
            taskC();
        }
    }
}
```




```
#pragma omp parallel sections
{
    #pragma omp section
    {
        taskA();
    }
    #pragma omp section
    {
        taskB();
    }
    #pragma omp section
    {
        taskC();
    }
}
```



Merging Directives (2)

```
#pragma omp parallel
{
    #pragma omp for
        for (i = 0; i < mmax; i++)
            /* body of loop */
}
```




```
#pragma omp parallel for
{
    for (i = 0; i < mmax; i++)
        /* body of loop */
}
```



Caution for Merging Directives (1)

- Each parallel directive forks threads
- Then, join threads after the parallel construct

```
#pragma omp parallel for  
for (i=0; i<n; ++i) {  
    ...  
}  
#pragma omp parallel for  
for (i=0; i<n; ++i) {  
    ...  
}
```




Caution for Merging Directives (2)

- Parallelize a loop using threads that are forked in advance

```
#pragma omp parallel num_threads(n)
{
    #pragma omp for
    for (i=0; i<n; ++i) {
        ...
    }
    #pragma omp for
    for (i=0; i<n; ++i) {
        ...
    }
}
```

fork threads for the following structured block

parallelize the following for loop using the pre-forked threads



Unnecessary fork&join is eliminated



Nesting `parallel` Directives

- Nested parallelism can be enabled using the `OMP_NESTED` environment variable.
 - If the `OMP_NESTED` environment variable is set to `TRUE`, nested parallelism is enabled.
 - In this case, each parallel directive creates a new team of threads.

