



Multicore Computing

Lecture11 - Parallel Sorting



남 범 석

bnam@skku.edu




Sorting Algorithms

- Process of arranging unordered items into order
- Internal versus external sorting
 - External sorting uses auxiliary storage
- Comparison-based
 - Compare pairs of elements and exchange
 - e.g.) quick sort, merge sort, etc
 - $O(n \log n)$
- Non-comparison-based
 - Use known properties of elements
 - e.g.) counting sort, radix sort, etc
 - $O(n)$

Counting Sort

- Idea: build a histogram of the keys and compute position in answer array for each element

- $A = [3, 5, 4, 1, 3, 4, 1, 4]$



- Make temp array B, and write values into position
 - $B = [1, 1, 3, 3, 4, 4, 4, 5]$
 - Cost = $O(\text{\#keys} + \text{size of histogram})$
 - What if histogram too large (eg all 32-bit ints? All words?)




Sorting on Parallel Computers

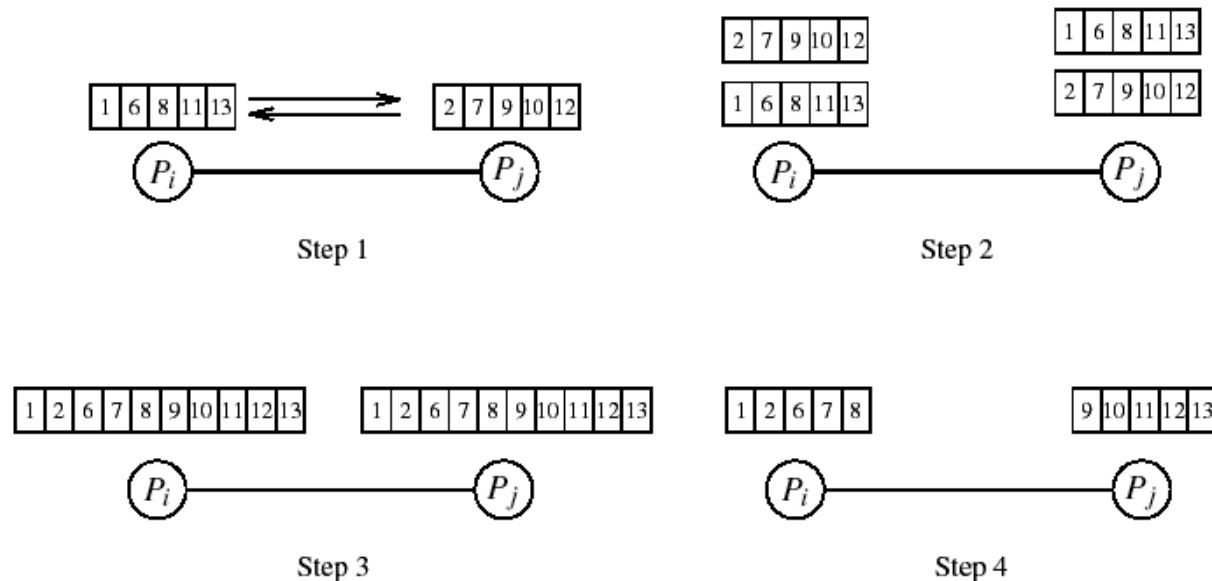
- $O(n \log n)$ optimal sequential sorting algorithm
- Comparison-based parallel sorting using n processors:
 - $O(n \log n) / n = O(\log n)$

Single vs. Multi Element Comparison

- One element per processor
 - Two numbers, say A and B , are compared between P_0 and P_1 .



- Multiple elements per processor



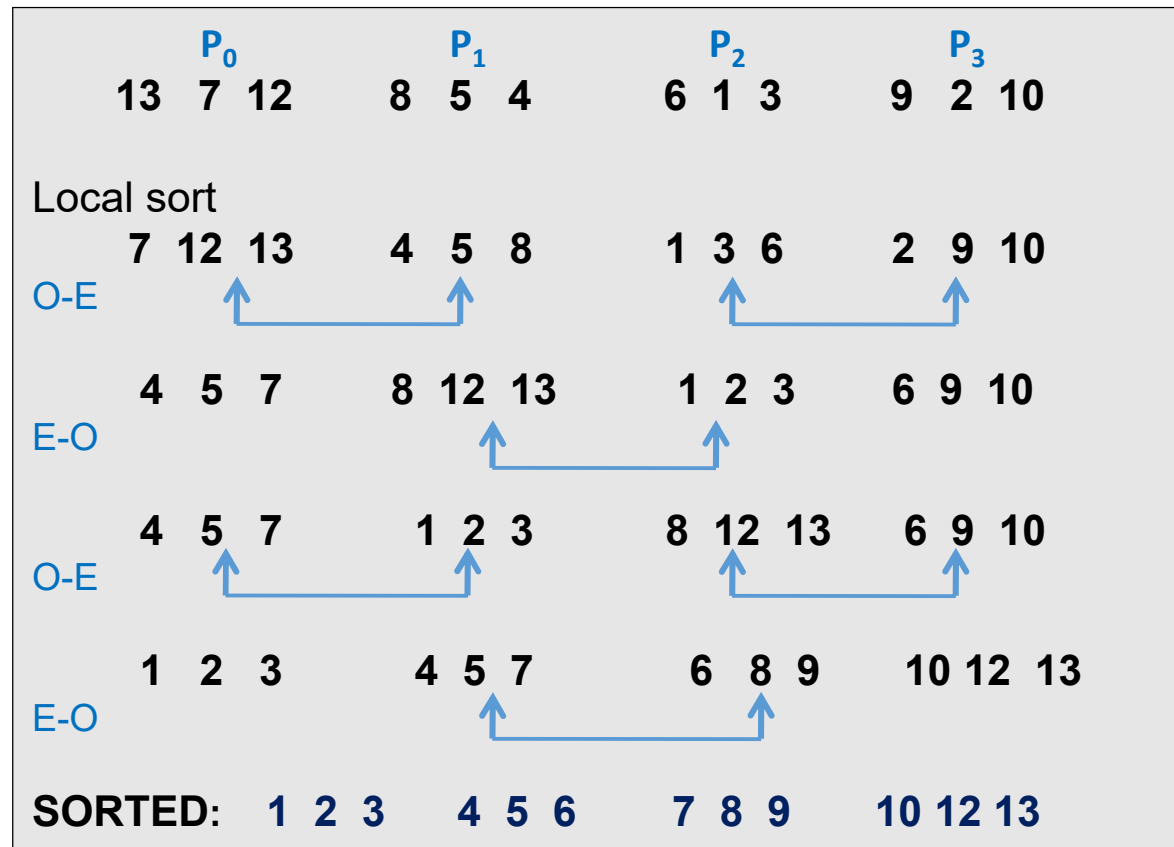
Odd-Even Transposition Sort

		P ₀	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆	P ₇				
Time ↓	Step 0	4	↔	2	7	↔	8	5	↔	1	3	↔	6
	1	2		4	↔	7	8	↔	1	5	↔	3	6
	2	2	↔	4	7	↔	1	8	↔	3	5	↔	6
	3	2		4	↔	1	7	↔	3	8	↔	5	6
	4	2	↔	1	4	↔	3	7	↔	5	8	↔	6
	5	1		2	↔	3	4	↔	5	7	↔	6	8
	6	1	↔	2	3	↔	4	5	↔	6	7	↔	8
	7	1	2	↔	3	4	↔	5	6	↔	7	8	

Parallel time complexity: $T_{par} = O(n)$ (for $P=n$)

Odd-Even Transposition Sort – Example ($N \gg P$)



- Each processor gets n/p numbers. First, sort n/p locally, then run odd-even transposition algorithm each time doing a merge-split for $2n/p$ numbers.



Time complexity: $T_{\text{par}} = (\text{Local Sort}) + (p \text{ merge-splits}) + (p \text{ exchanges})$


$$T_{\text{par}} = (n/p)\log(n/p) + p \cdot (n/p) + p \cdot (n/p) = (n/p)\log(n/p) + 2n$$

Parallelizing Merge-Sort



Process allocation

Parallelizing Merge-Sort



Sequential :

$$T_{seq} = 1 * n + 2 * \frac{n}{2} + 2^2 * \frac{n}{2^2} + \dots + 2^{\log n} * \frac{n}{2^{\log n}}$$

$$T_{seq} = O(n \log n)$$

Parallel :

$$T_{par} = 2 \left(\frac{n}{2^0} + \frac{n}{2^1} + \frac{n}{2^2} + \dots + \frac{n}{2^k} \dots + 2 + 1 \right)$$

$$= 2n \left(2^0 + 2^{-1} + 2^{-2} + \dots + 2^{-\log n} \right)$$

$$T_{par} = O(4n)$$

Bitonic Sort

- Create a *bitonic sequence* then sort the sequence
- *Bitonic sequence*
 - sequence of elements $\langle a_0, a_1, \dots, a_{n-1} \rangle$ is bitonic if
 - $\exists i, 0 \leq i \leq n-1$ s.t $\langle a_0, a_1, \dots, a_i \rangle$ is monotonically increasing
 - And $\langle a_i, a_{i+1}, \dots, a_{n-1} \rangle$ is monotonically decreasing

Or



- $\exists i, 0 \leq i \leq n-1$ s.t $\langle a_0, a_1, \dots, a_i \rangle$ is monotonically decreasing
- And $\langle a_i, a_{i+1}, \dots, a_{n-1} \rangle$ is monotonically increasing



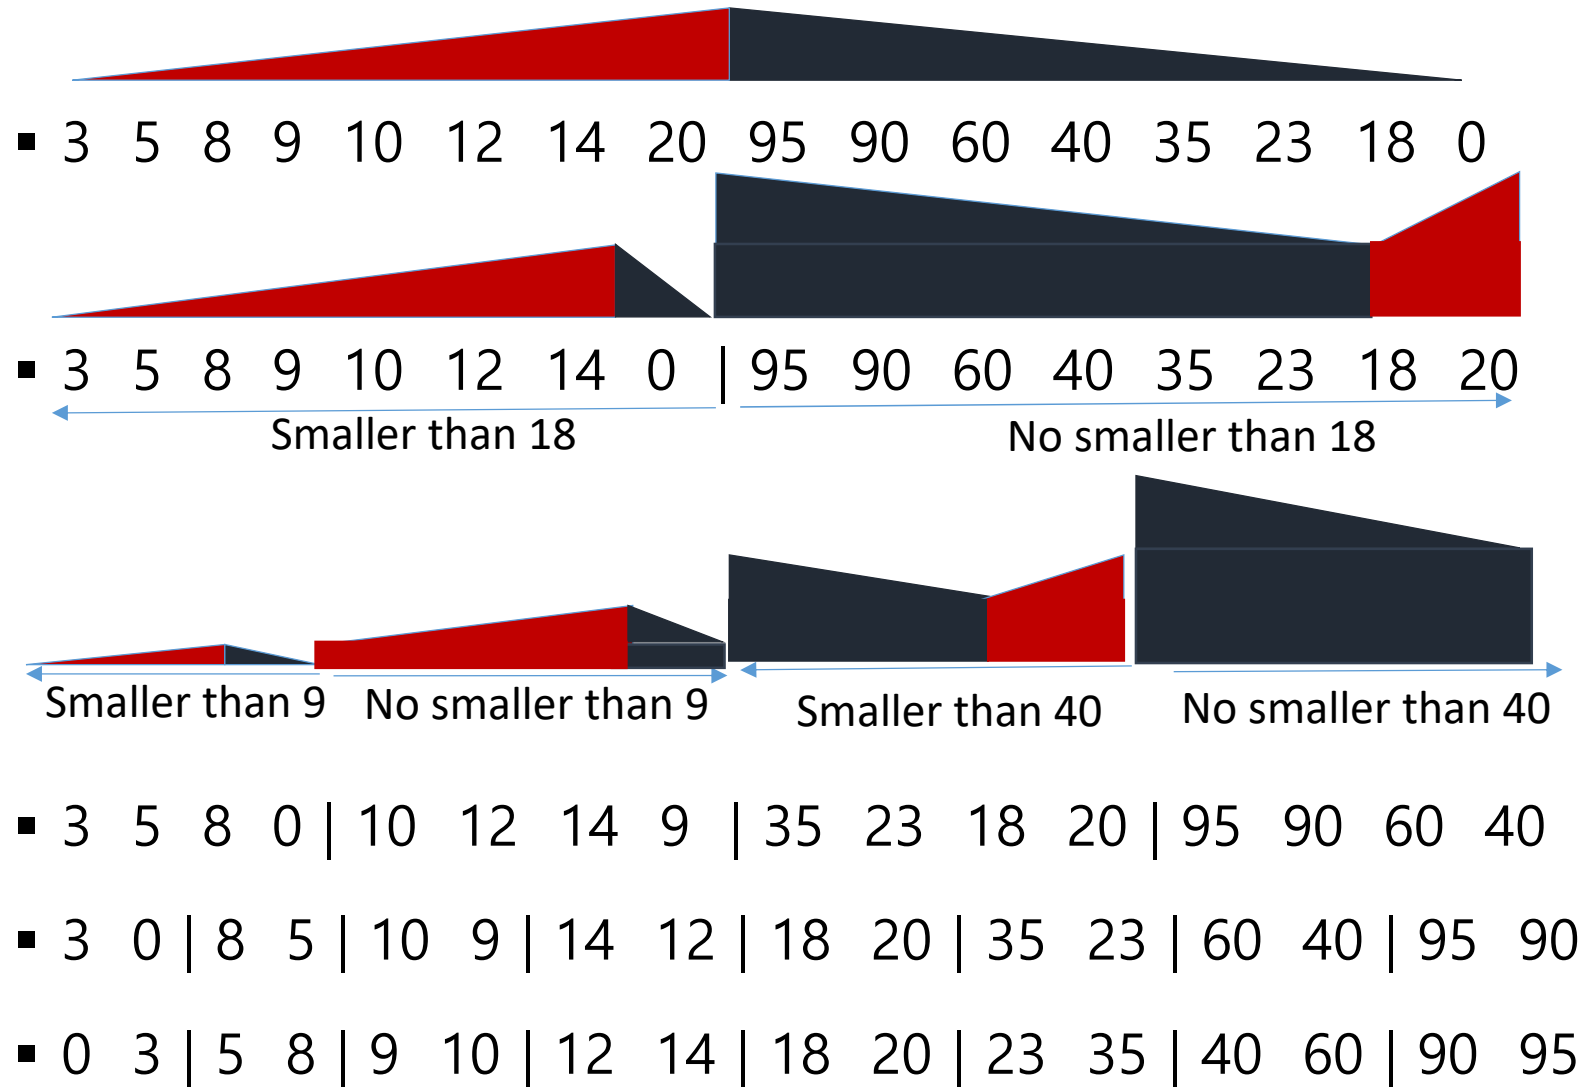
- For example,
 - $\langle 1, 4, 6, 8, 3, 2 \rangle$ and $\langle 9, 8, 3, 2, 4, 6 \rangle$ are both bitonic

Bitonic Merge: Sorting Bitonic Sequence




- Compare the first element of s_1 with the first element of s_2 , then the second element of s_1 with second element of s_2 and so on.
- We exchange elements if an element of s_2 is smaller.
- After above compare and exchange steps, we get two bitonic sequences of length $n/2$ such that all elements in first bitonic sequence are smaller than all elements of second bitonic sequence.
- We repeat the same process within two bitonic sequences and we get four bitonic sequences of length $n/4$ such that all elements of leftmost bitonic sequence are smaller and all elements of rightmost.
- If all these bitonic sequence are sorted and every bitonic sequence has one element, we get the sorted array.

Bitonic Merge - Example




Sorting Network

- Network of comparators designed for sorting
- Comparator : two inputs x and y ; two outputs x' and y'
 - Two types
 - increasing (denoted \oplus): $x' = \min(x,y)$ and $y' = \max(x,y)$



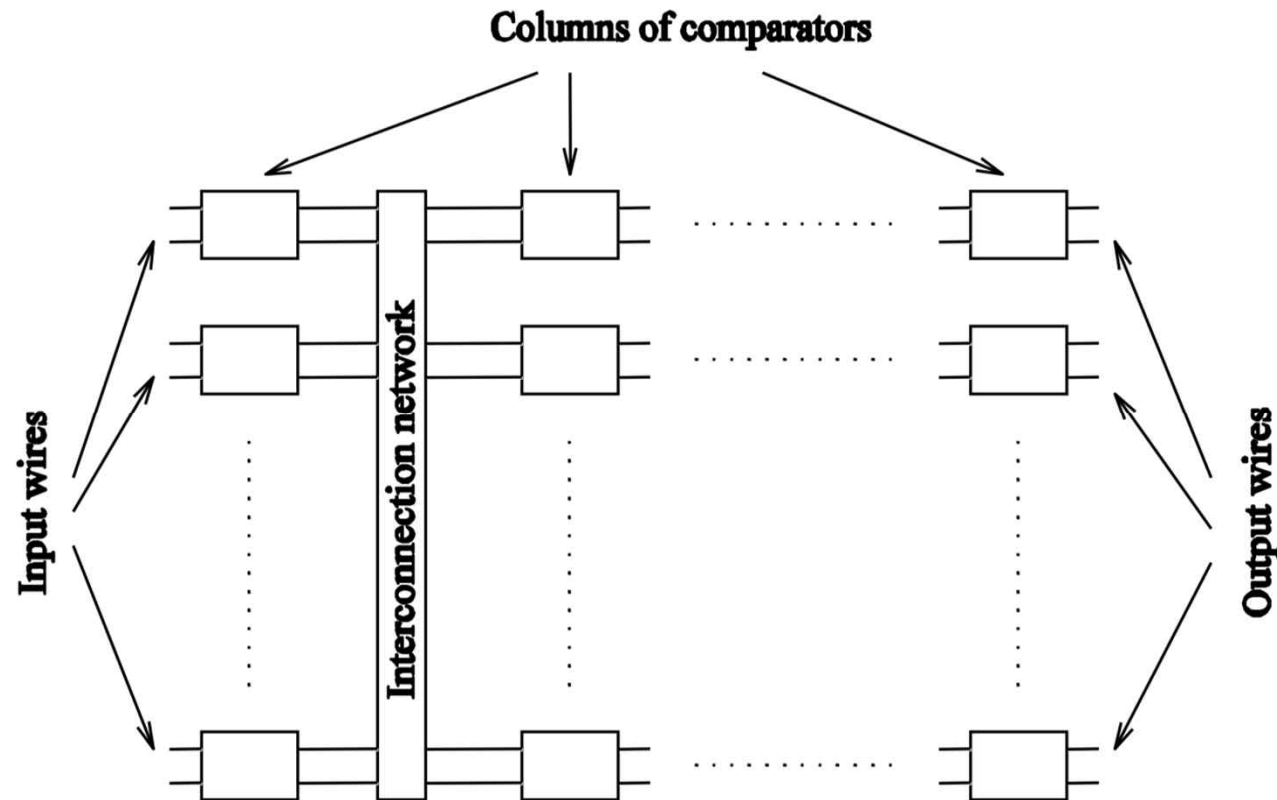
- decreasing (denoted \ominus) : $x' = \max(x,y)$ and $y' = \min(x,y)$



- Sorting network speed is proportional to its depth

Sorting Network

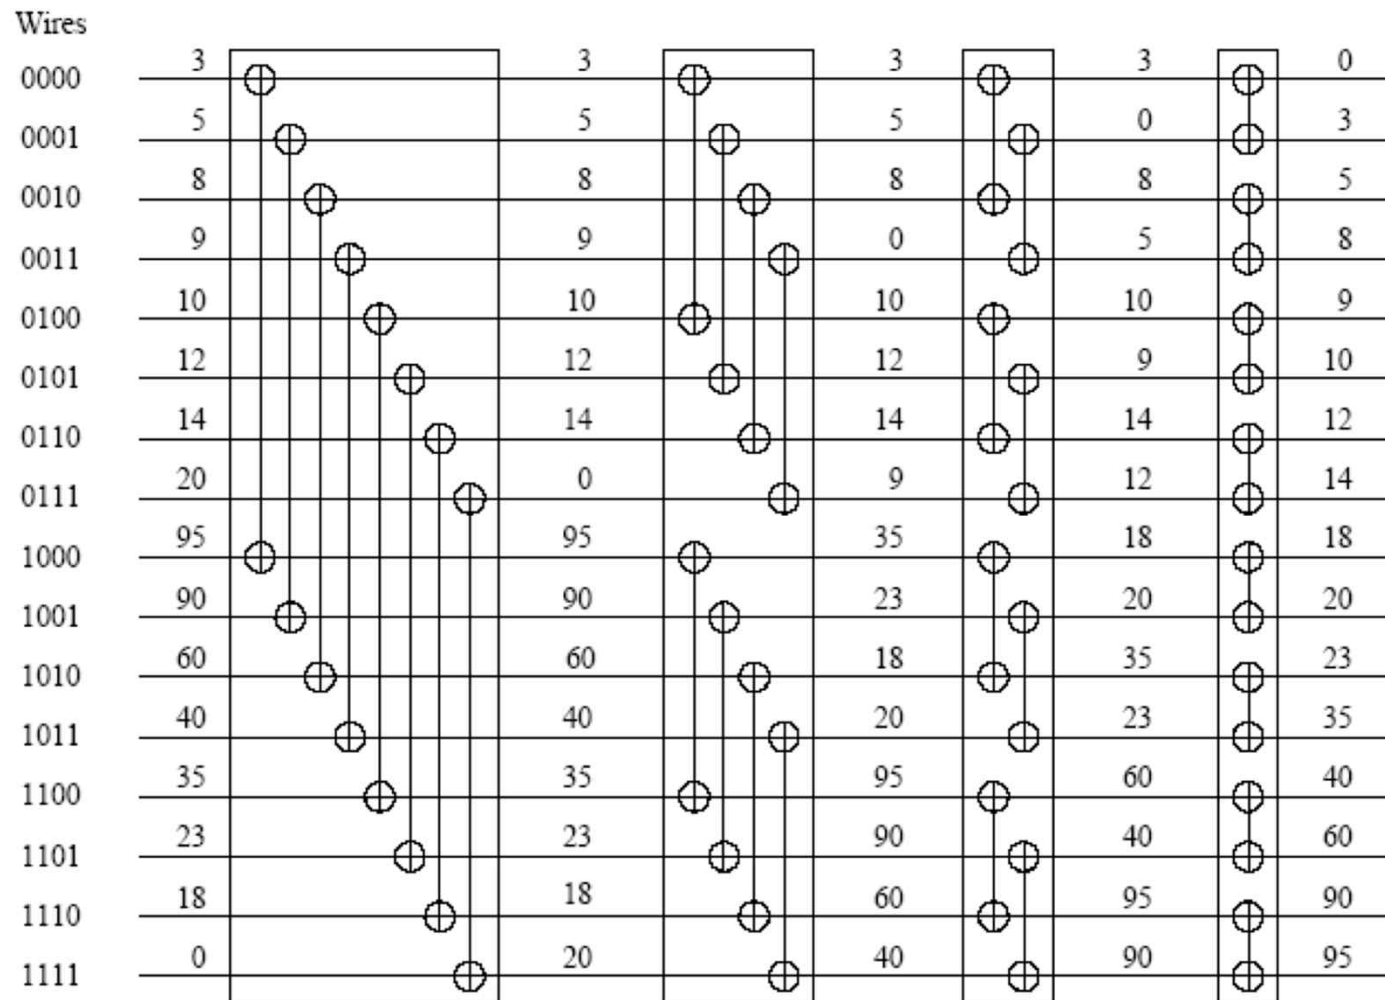
- Network structure: a series of columns
- Each column consists of a vector of parallel comparators



Bitonic Merge Network

- Input: Bitonic Sequence

\oplus BM[16]






Bitonic Sort

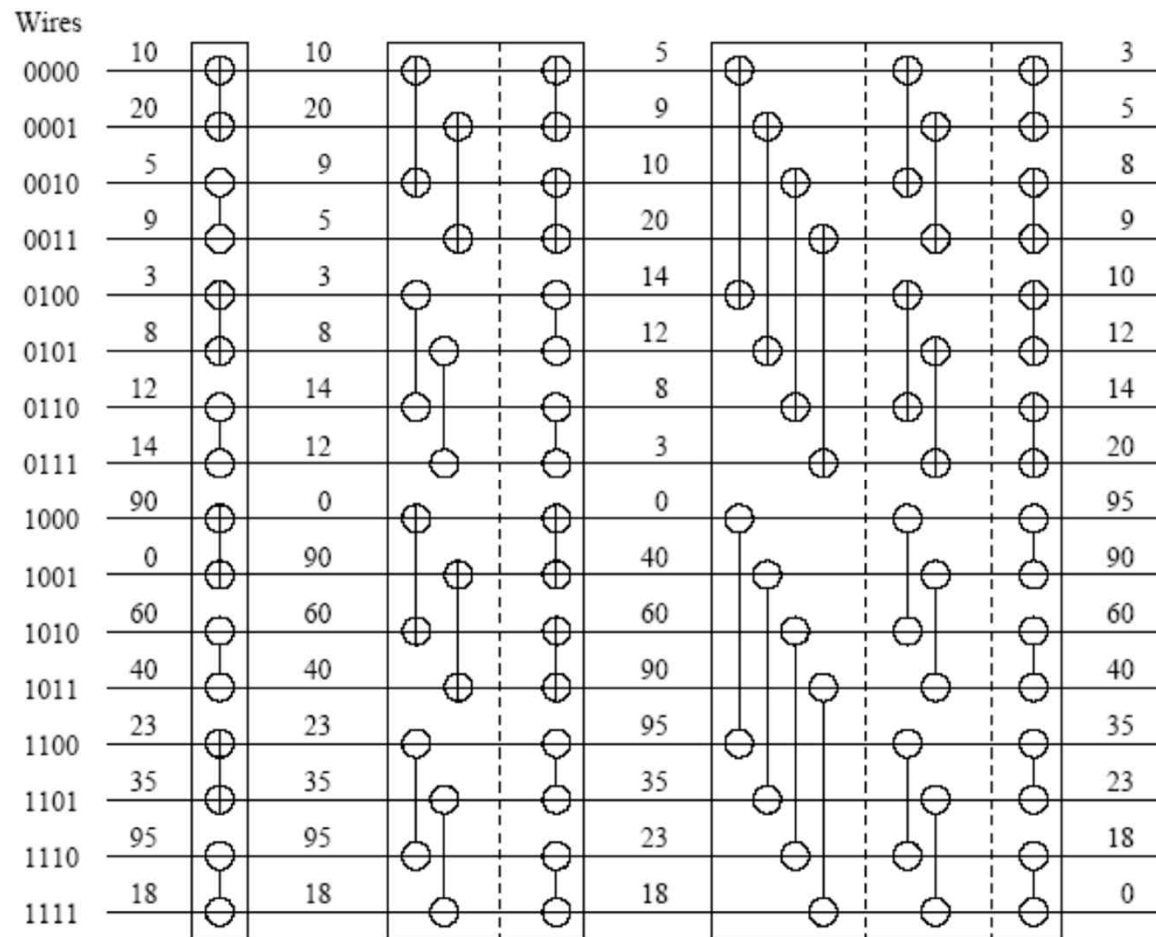
- Q: How to sort an unsorted sequence ?
- Two steps
 - Build a bitonic sequence
 - Sort it using a bitonic merging network

Building a Bitonic Sequence

- Build a single bitonic sequence from the given sequence
 - any sequence of length 2 is a bitonic sequence.
 - build bitonic sequence of length 4
 - sort first two elements using $\oplus \text{BM}[2]$
 - sort next two using $\ominus \text{BM}[2]$
- Repeatedly merge to generate larger bitonic sequences
 - $\oplus \text{BM}[k]$ & $\ominus \text{BM}[k]$: bitonic merging networks of size k




Building a Bitonic Sequence



- Input: sequence of 16 unordered numbers
- Output: a bitonic sequence of 16 numbers

Bitonic Sort, $n = 16$




- First 3 stages create bitonic sequence input to stage 4
- Last stage ($\oplus \text{BM}[16]$) yields sorted sequence

Bitonic Sort - Example




- Convert the following sequence to bitonic sequence:
- 3, 7, 4, 8, 6, 2, 1, 5
- [Step 1] Consider each 2-consecutive elements as bitonic sequence and apply bitonic sort on each 2- pair elements.

$$\begin{array}{l} a \\ b \end{array} \begin{array}{l} \searrow \\ \nearrow \end{array} \Rightarrow \begin{array}{l} \min(a,b) \\ \max(a,b) \end{array} \quad \begin{array}{l} a \\ b \end{array} \begin{array}{l} \nearrow \\ \searrow \end{array} \Rightarrow \begin{array}{l} \max(a,b) \\ \min(a,b) \end{array}$$




Bitonic Sort - Example

- Convert the following sequence to bitonic sequence:
- 3, 7, 4, 8, 6, 2, 1, 5
- [Step 2] In next step, take two 4 element bitonic sequences and so on.



Bitonic Sort - Example

- Convert the following sequence to bitonic sequence:
- 3, 7, 4, 8, 6, 2, 1, 5
- [Step 3] In next step, take one 8 element bitonic sequence and do the bitonic merge.



Bitonic Sort – Sequential C code

```
void bitonicSort(int a[],int low, int cnt, int dir)
{
    if (cnt>1)
    {
        int k = cnt/2;

        // sort in ascending order since dir here is 1
        bitonicSort(a, low, k, 1);

        // sort in descending order since dir here is 0
        bitonicSort(a, low+k, k, 0);

        // Will merge whole sequence in ascending order
        // since dir=1.
        bitonicMerge(a,low, cnt, dir);
    }
}
```

```
void bitonicMerge(int a[], int low, int cnt, int dir)
{
    if (cnt>1)
    {
        int k = cnt/2;
        for (int i=low; i<low+k; i++)
            compAndSwap(a, i, i+k, dir);
        bitonicMerge(a, low, k, dir);
        bitonicMerge(a, low+k, k, dir);
    }
}
```




Bitonic Sort

- A lot more comparisons than other sorting algorithms
- But
 - Compares elements in predefined sequence
 - Sequence of comparison doesn't depend on data
- Therefore, more suitable for parallel computing


Shear Sort

- The layout of a sorted sequence on a mesh could be row by row or snakelike:




Shear Sort


- Alternate row and column sorting until list is fully sorted.
- Alternate row directions to get snake-like sorting:




(a) Original placement of numbers




(b) Phase 1 — Row sort




(c) Phase 2 — Column sort



(d) Phase 3 — Row sort



(e) Phase 4 — Column sort



(f) Final phase — Row sort



Shear Sort - Complexity

- On a $n \times n$ mesh, it takes $2 \log(n)$ phases to sort n^2 numbers.

$$T_{par}^{shearsort} = O(n \log n) \quad \text{on a } n \times n \text{ mesh}$$

- Sequential sorting takes $T = O(n^2 \log n)$. Therefore,

$$Speedup_{shearsort} = \frac{T_{seq}}{T_{par}} = O(n) \quad (\text{for } P = n^2)$$

$$\text{However, efficiency} = \frac{1}{n}$$

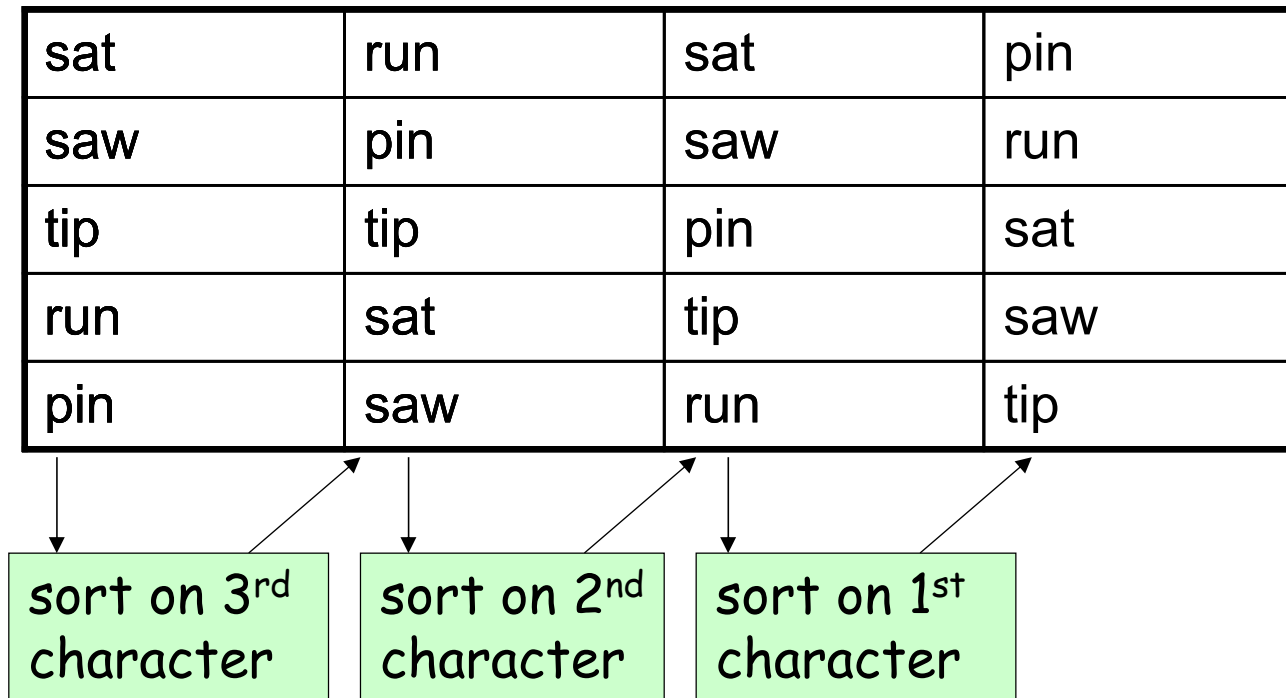


Radix Sort

- Assume numbers are in positional digit representation.
 - The digits represent values
 - The position of each digit indicates their relative weighting
 - E.g., binary or decimal numbers
- Radix sort starts at the least significant digit and sorts the numbers according to their digits.
- The sequence is then sorted according to the next least significant digit and so on until the most significant digit, after which the sequence is sorted.

Radix Sort: Separate Key Into Parts


- Divide keys into parts, e.g., by digits (radix)
- Using counting sort on these each radix:
 - Start with least-significant



- Cost = $O(\text{\#keys} * \text{\#characters})$

Radix Sort

- Example: Radix Sort using binary digits



- Ok, it's an easy algorithm.
- But how do we know where to put each element, i.e., new location, if it has to move?
 - Build histograms



Parallel Radix Sort

- Overview:
- 1. Count occurrences of each digit/process pair
 - **Histogram[p][d]**: Occurrences of digit d on process p
 - Local computation
- 2. Find global block offset for each digit/process pair
 - **offset[p][d]**: Block offset for digit d on proc p
 - Prefix sum of Histogram[][] in column-major order
- 3. Find destination index for each input element
 - **index[p][k]**: Destination for element k on p
- 4. Apply the input-to-output permutation
 - $\text{output}[\text{index}[k]] = \text{input}[k]$
 - Permutation / all-to-all

Parallel Radix Sort

- Build a histogram for each bit

Element #	0	1	2	3	4
Value:	7	14	4	1	6
Binary:	0111	1110	0100	0001	0110
bit 0:	1	0	0	1	0

- →
Histogram: Zero bits One bit
 3 2
- prefix sum on these histogram values

- →
Histogram: Zero bits One bit
 3 2

Prefix Sum: 0 3

- Determine the relative offset using prefix sum

Element #	0	1	2	3	4
Value:	7	14	4	1	6
Binary:	0111	1110	0100	0001	0110
bit 0:	1	0	0	1	0
offset:	0	0	1	1	2

Parallel Radix Sort

- Determine the location to move

Element #	0	1	2	3	4
Value:	7	14	4	1	6
Binary:	0111	1110	0100	0001	0110
bit 0:	1	0	0	1	0
offset:	0	0	1	1	2
index:	3	0	1	4	2

Element #	0	1	2	3	4
Value:	14	4	6	7	1
Binary:	1110	0100	0110	0111	0001

```
graph TD
    subgraph Initial [Initial Array]
        direction TB
        E0[Element # 0]
        V0[Value: 7]
        B0[Binary: 0111]
        B0_0[bit 0: 1]
        O0[offset: 0]
        I0[index: 3]
    end
    subgraph Sorted [Sorted Array]
        direction TB
        E1[Element # 0]
        V1[Value: 14]
        B1[Binary: 1110]
    end
    I0 --> E1
    I1[Element # 1] --> E0
    I2[Element # 2] --> E2[Element # 2]
    I3[Element # 3] --> E3[Element # 3]
    I4[Element # 4] --> E4[Element # 4]
```

- Repeat for the next bit
- Q: Complexity?