



Multicore Computing

Lecture14 – Lock-Free Hashing




남 범 석

bnam@skku.edu



Background: Hashing



5 Items

$$h(k) = k \bmod 4$$






Lock-Free Hashing: No Brainer?

- Concurrent Hashing: Add(x), Remove(x), Contains(x)
 - Simple
 - Lock-free
- However, we don't know how to resize ...



Background: Hashing



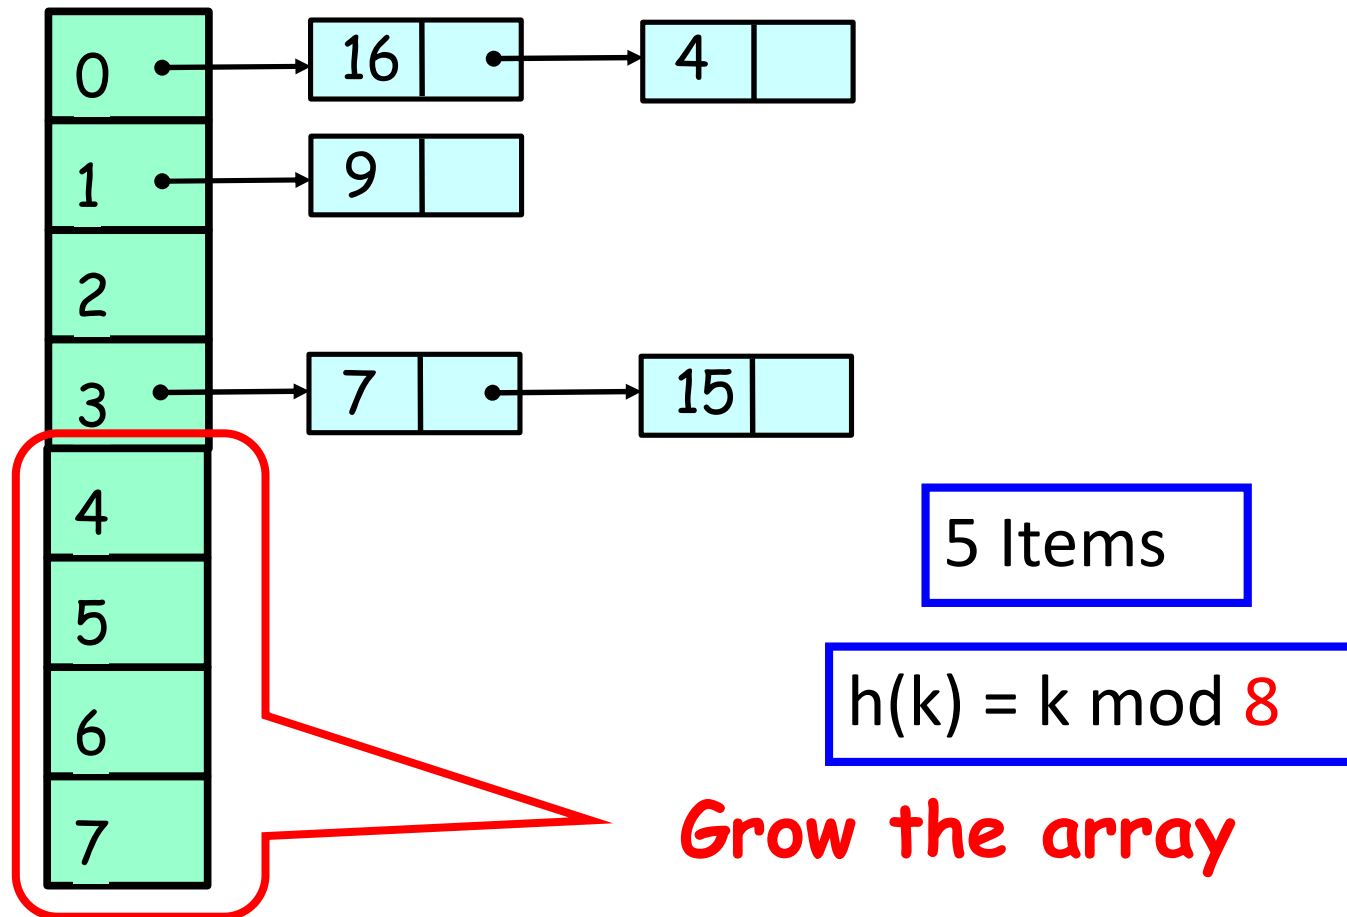
Problem:
buckets getting too long

5 Items

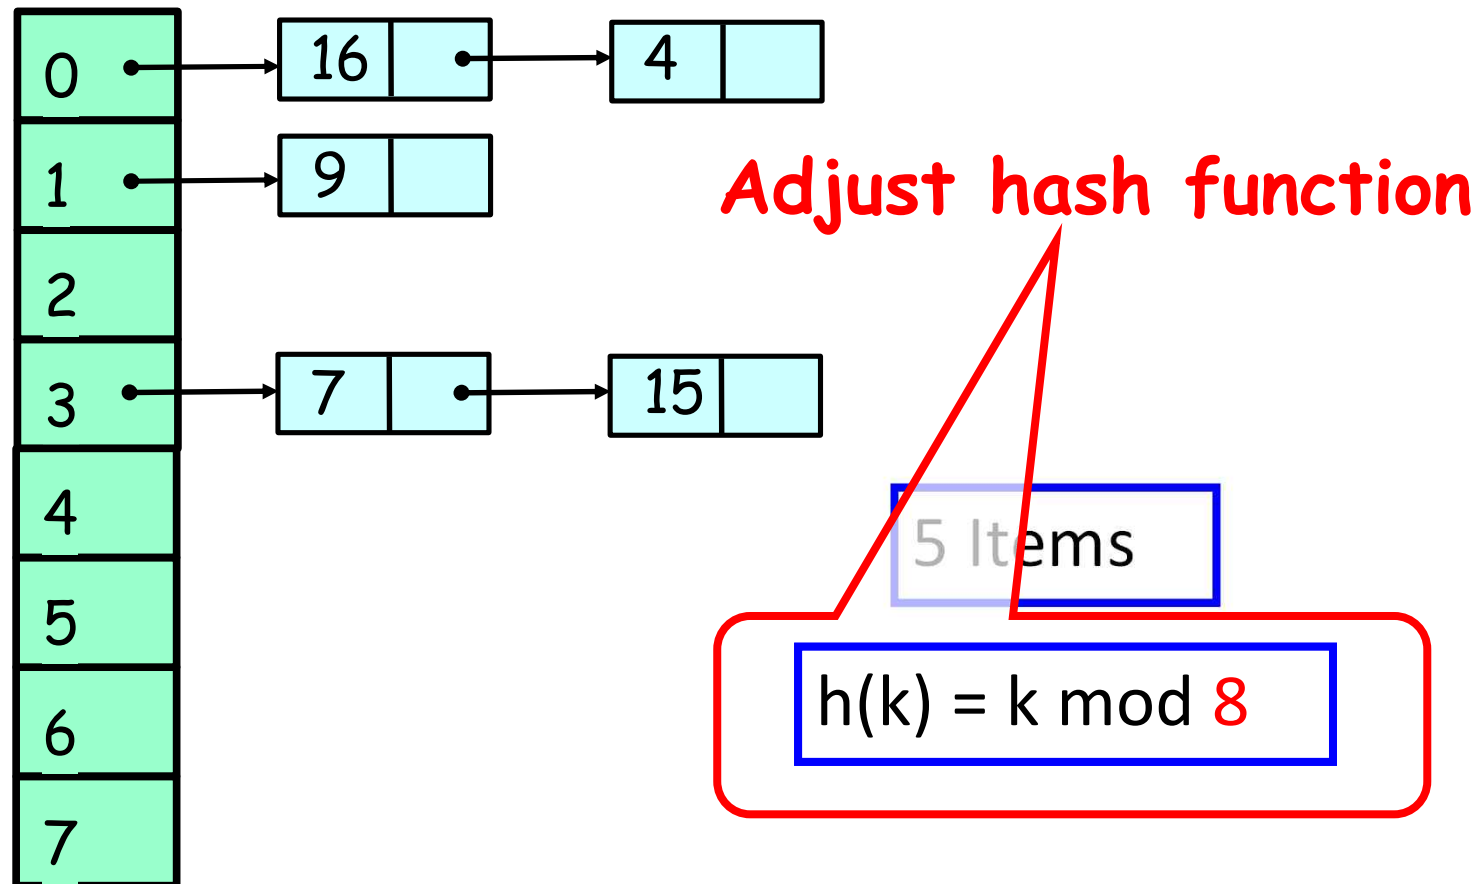
$$h(k) = k \bmod 4$$



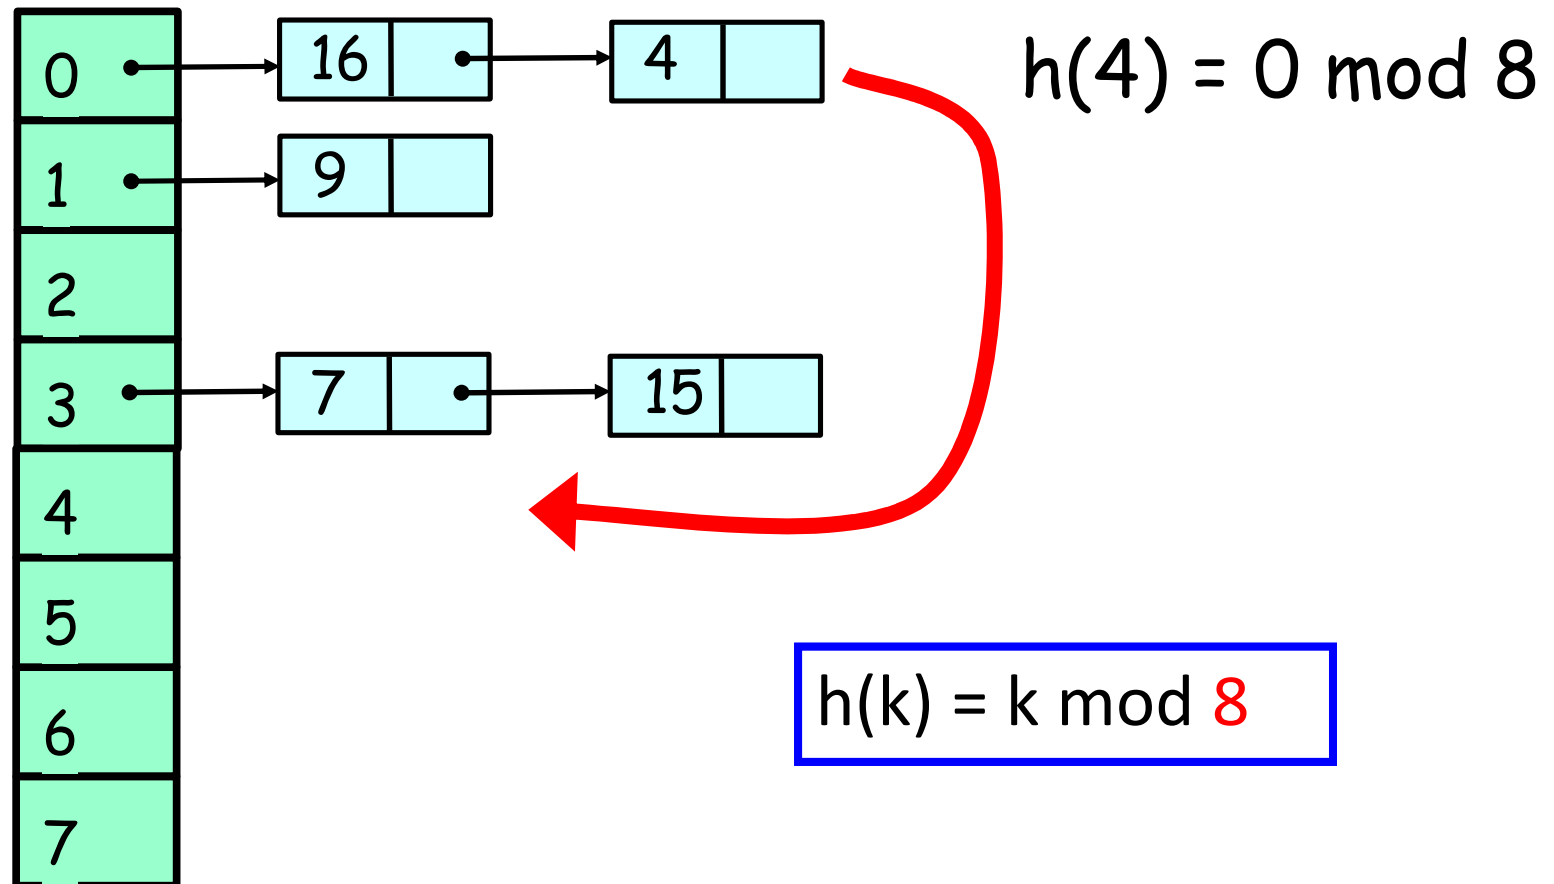
Background: Hashing



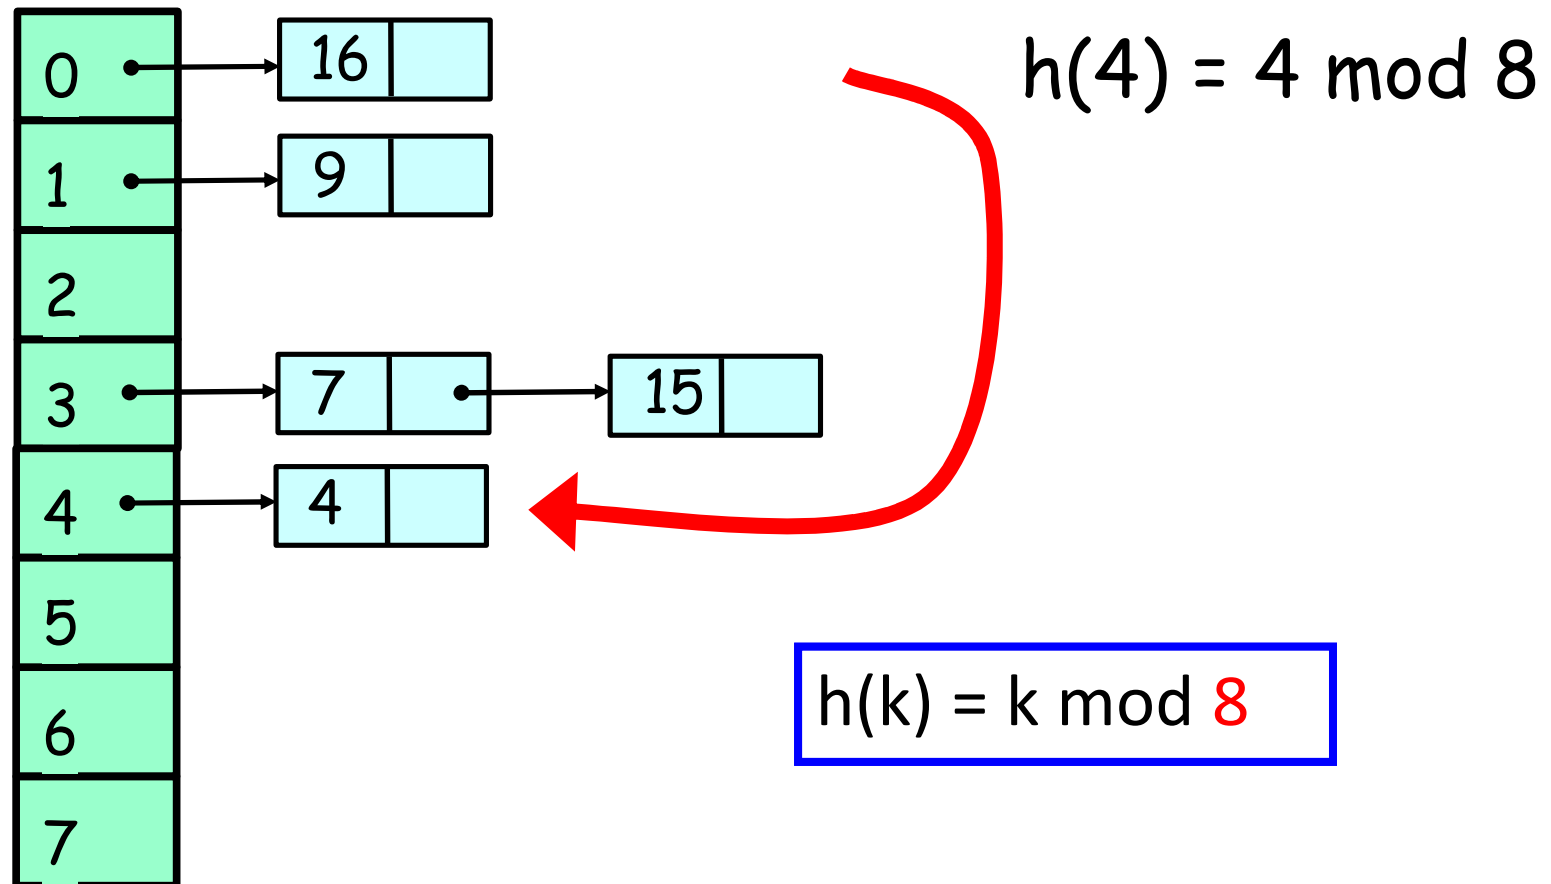
Background: Hashing



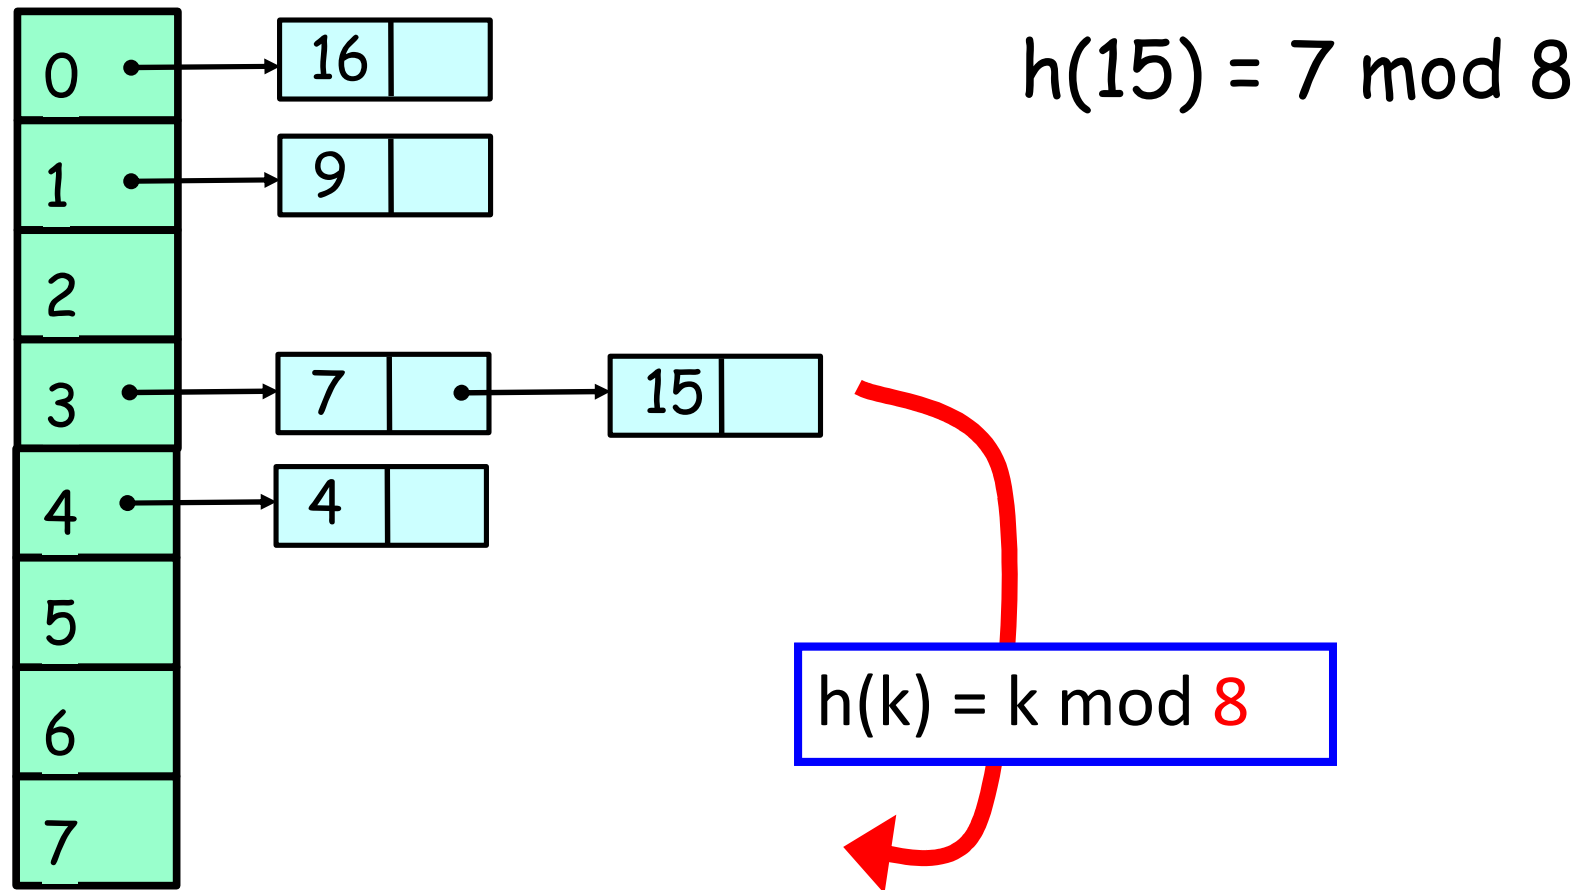
Background: Re-Hashing (Resizing)



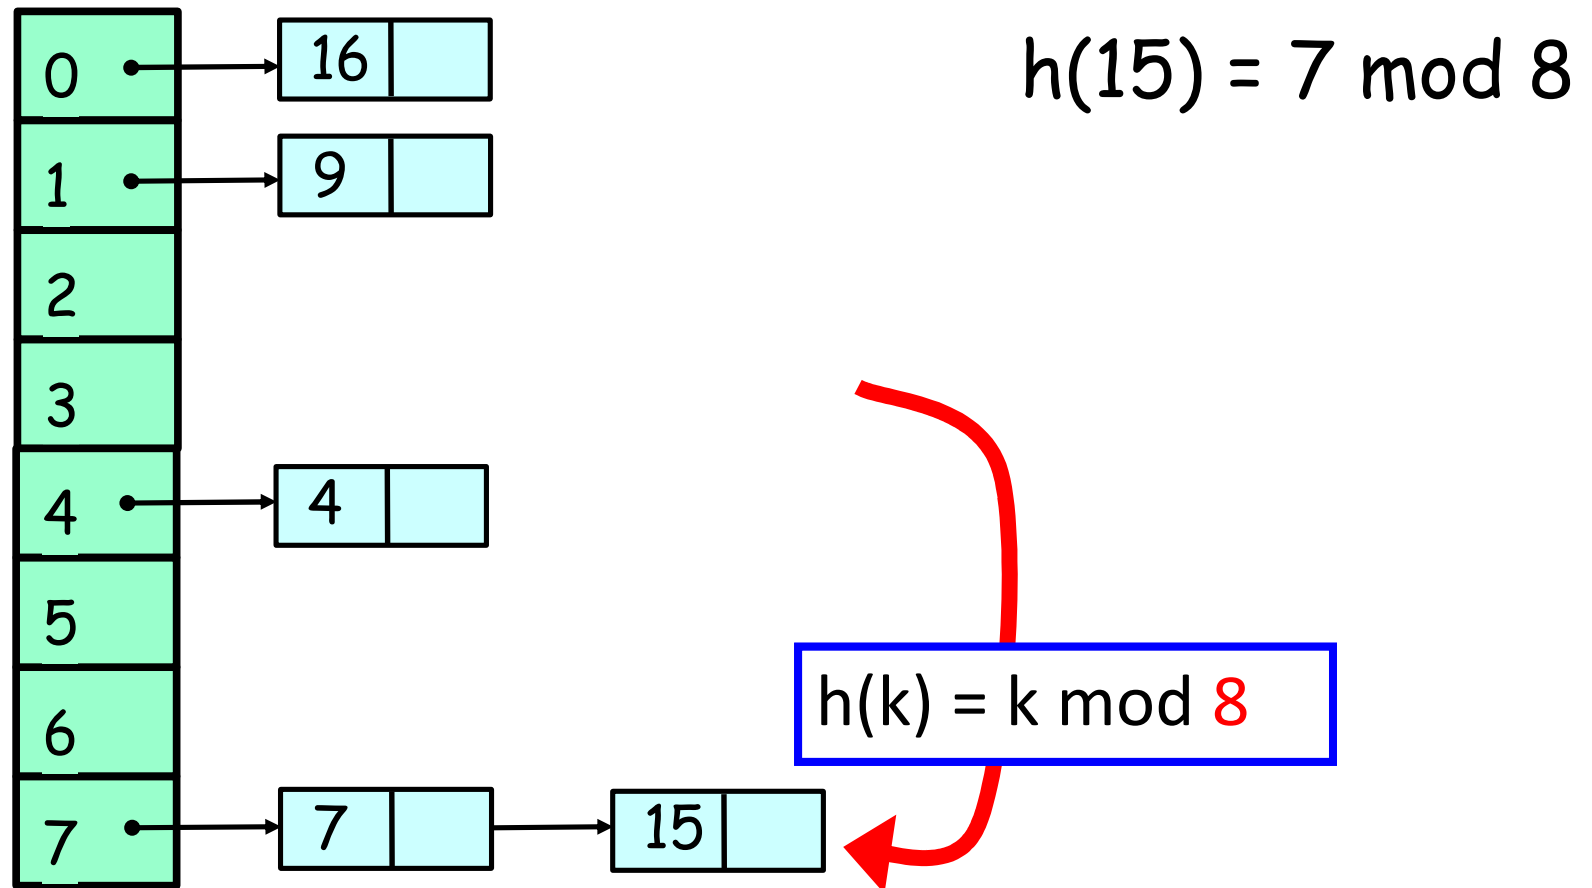
Background: Re-Hashing (Resizing)



Background: Re-Hashing (Resizing)



Background: Re-Hashing (Resizing)





Is Resizing Necessary?

- Constant-time method calls require
 - Constant-length buckets
 - Table size proportional to set size
 - As set grows, must be able to resize






Coarse-Grained Locking

- Good parts
 - Simple
 - Hard to mess up
- Bad parts
 - Sequential bottleneck



Fine-grained Locking




Each lock associated with one bucket




Fine-grained Locking: Resize This

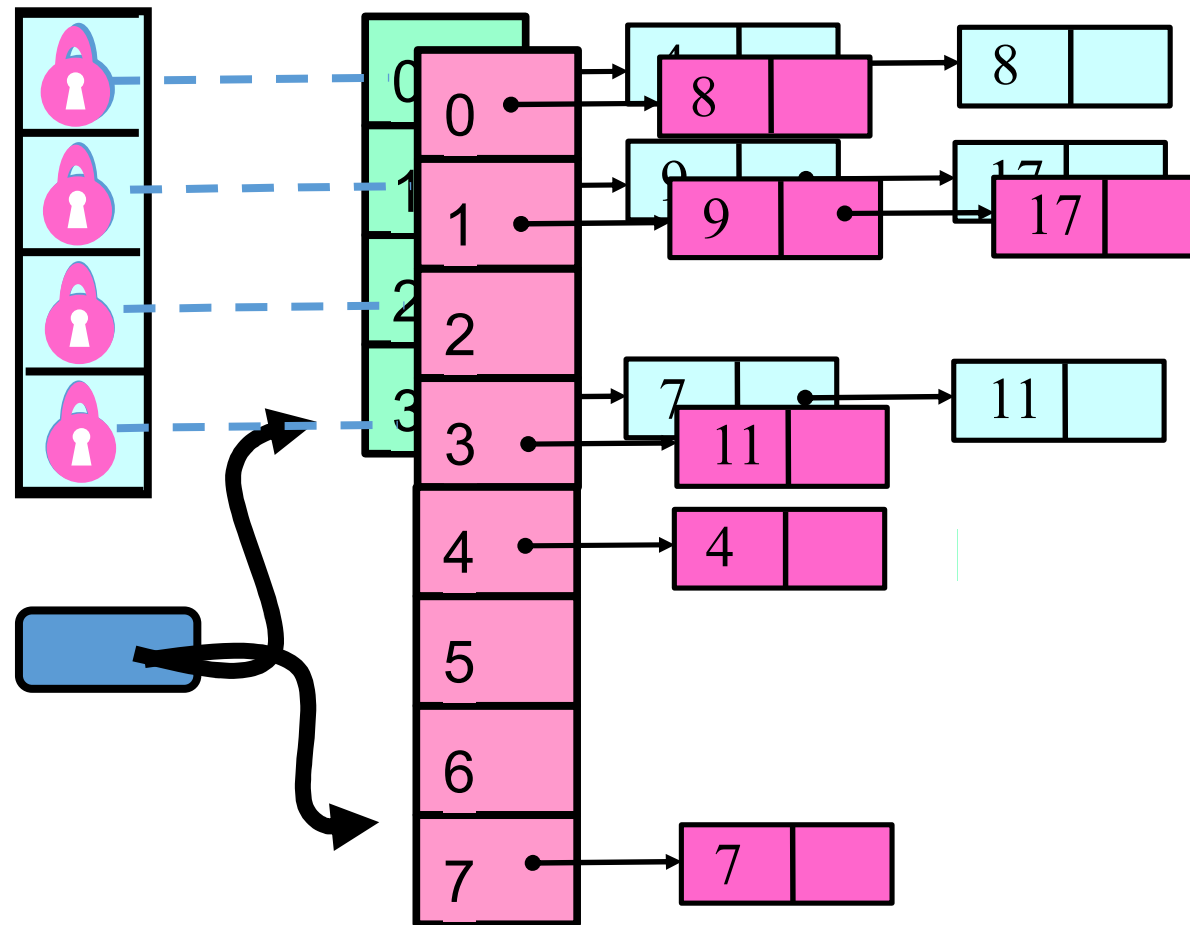
- Make sure table reference didn't change between resize decision and lock acquisition



Fine-grained Locking: Resize This

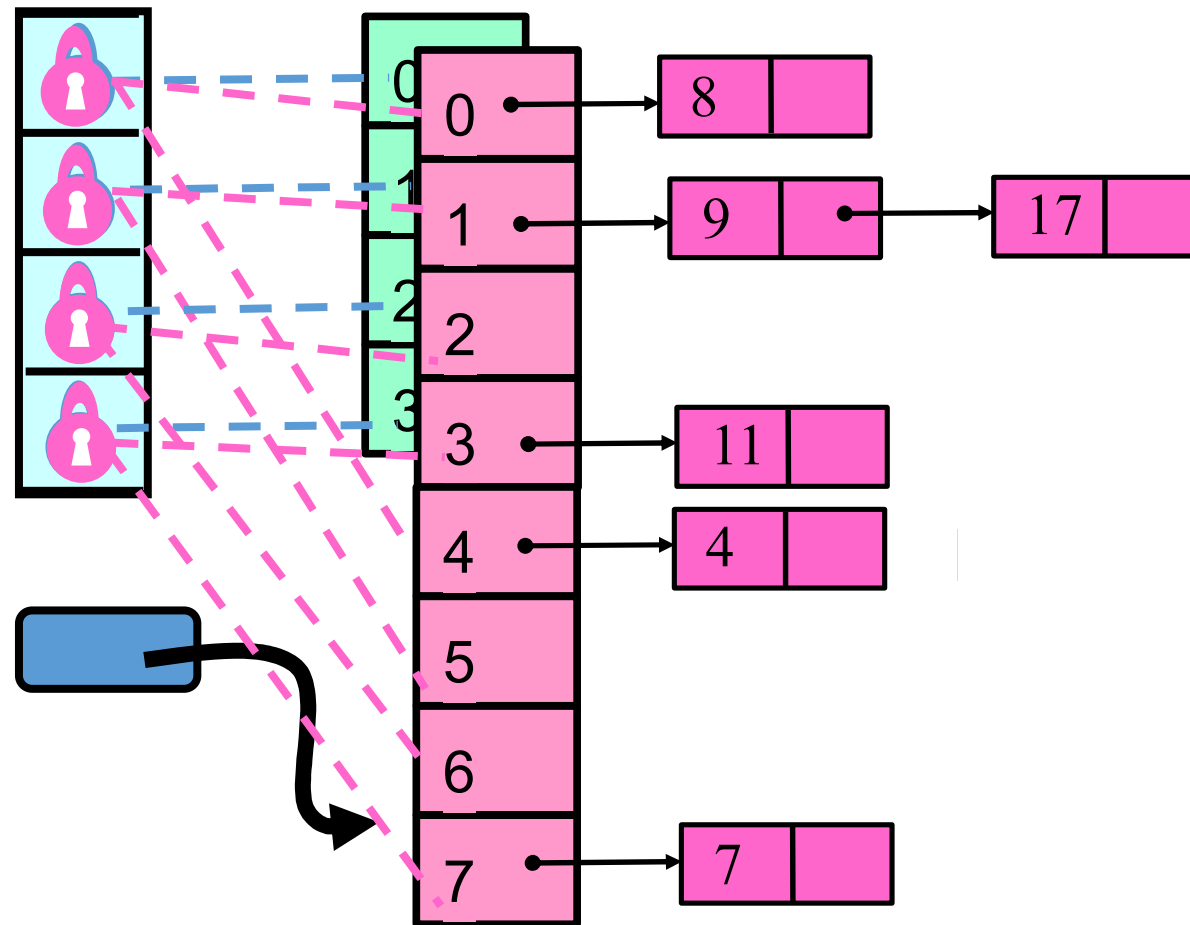


Fine-grained Locking: Resize This



Fine-grained Locking: Resize This

Striped Locks: each lock now associated with two buckets





Fine-grained Locking: Observations

- We grow the table, but not locks
 - Resizing lock array is tricky ...
- We use sequential lists
 - Not LockFreeList lists
 - If we're locking anyway, why pay?





Fine-grained Locking:

- We can resize the table
- But not the locks
- Debatable whether method calls are constant-time in presence of contention ...






Stop The World Resizing


- Resizing stops all concurrent operations
- What about an incremental resize?
- Must avoid locking the table
- A lock-free table + incremental resizing?



Lock-Free Resizing Problem




Lock-Free Resizing Problem



Need to extend table




Lock-Free Resizing Problem




Need to extend table



Lock-Free Resizing Problem



Lock-Free Resizing Problem




to remove and
then add even a
single item single
location CAS
not enough

We need a new idea...




Don't move the items


- Move the buckets instead
- Keep all items in a single lock-free list
- Buckets become "shortcut pointers" into the list




Recursive Split Ordering




Recursive Split Ordering



Recursive Split Ordering




Recursive Split Ordering




List entries sorted in order that allows recursive splitting. How?



Recursive Split Ordering




Recursive Split Ordering



LSB = Least significant Bit



Recursive Split Ordering



Split-Order

- If the table size is 2^i ,
 - Bucket b contains keys k
 - $k \bmod 2^i = b$
 - bucket index consists of key's i LSBs



When Table Splits

- Some keys stay
 - $b = k \bmod(2^i)$
- Some move
 - $b + 2^i = k \bmod(2^i)$
- Determined by $(i+1)$ st bit
 - Counting backwards
- Key must be accessible from both
 - Keys that will move must come later



A Bit of Magic


Real keys:

0		4		2		6		1		5		3		7	
---	--	---	--	---	--	---	--	---	--	---	--	---	--	---	--




A Bit of Magic

Real keys:



Real key 1 is in
the 4th location

Split-order:



A Bit of Magic

Real keys:

0	4	2	6	1	5	3	7
000	100	010	110	001	101	011	111

Real key 1 is in 4th location

Split-order:

0	1	2	3	4	5	6	7
000	001	010	011	100	101	110	111



A Bit of Magic

Real keys:

000 100 010 110 001 101 011 111

Split-order:

000 001 010 011 100 101 110 111



A Bit of Magic

Real keys:

000 100 010 110 001 101 011 111

Split-order:


000 001 010 011 100 101 110 111

**Just reverse the order of
the key bits**




Split Ordered Hashing


Order according to reversed bits



Parent Always Provides a Short Cut




Sentinel Nodes



Problem: how to remove a node pointed by 2 sources using CAS



Sentinel Nodes



Solution: use a Sentinel node for each bucket





Sentinel vs Regular Keys

- Want sentinel key for i ordered
 - before all keys that hash to bucket i
 - after all keys that hash to bucket $(i-1)$






Splitting a Bucket


- We can now split a bucket
- In a lock-free manner
- Using two CAS() calls ...



Initialization of Buckets



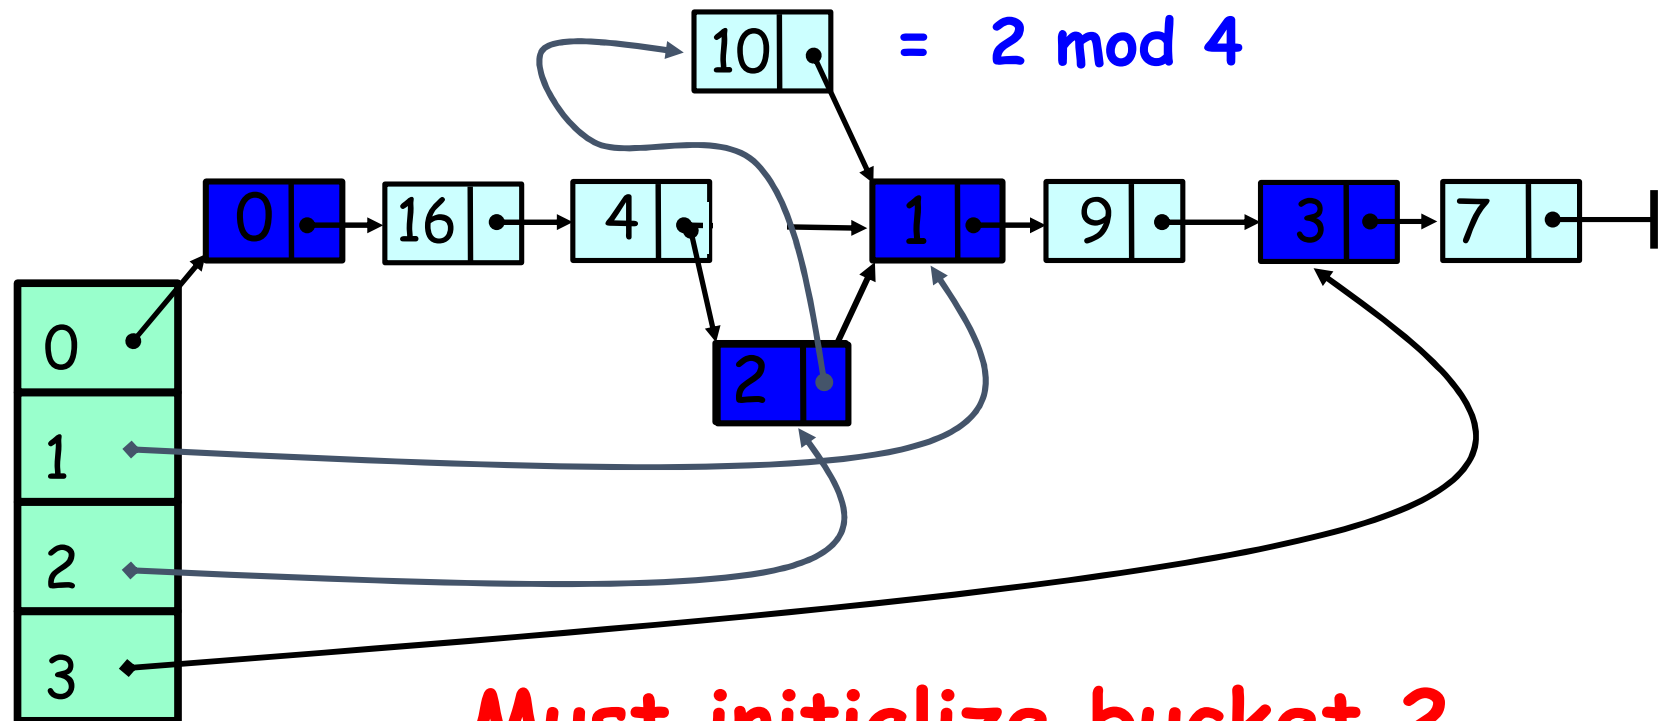
Initialization of Buckets



Need to initialize bucket 3 to split bucket 1




Adding 10



**Must initialize bucket 2
Then can add 10**



Recursive Initialization





Main List

- Lock-Free List from earlier class
- With some minor variations





Summary

- Concurrent resizing is tricky
- Lock-based
 - Fine-grained
 - Read/write locks
 - Optimistic
- Lock-free
 - Builds on lock-free list