



Multicore Computing

Lecture23 – CAP Theorem




남 범 석


bnam@skku.edu





Distributed Consensus



Let's have coffee



Let's have tea




Agreement in Faulty Systems

- Reaching agreement using E-mail becomes even harder.
 - Alice \rightarrow Bob : Let's meet at noon on the 2nd floor
 - Alice \leftarrow Bob : Ok!
 - Alice (What if Bob doesn't know that I received his message?)
 - Alice \rightarrow Bob : I received your message, so it's ok.
 - Bob (What if Alice doesn't know that I received her message?)
 - ...




Agreement in Faulty Systems

- Byzantine Agreement (Pease, Shostak, Lamport, 1980)
 - The Byzantine generals announce their troop strengths.
 - There are loyal generals and traitors as well.
 - Goal: Each troop learns the true strengths sent by each of the loyal generals.



Agreement in Faulty Systems ($n=3$)



**There is no solution for
3 Generals, 1 Traitor.**

CAP Theorem

- It is impossible for a web service to provide following *three guarantees at the same time*:
 - **Consistency**
 - **Availability**
 - **Partition-tolerance**
- Prof. Eric Brewer (U.C Berkeley and Google)






CAP Theorem

- Consistency:
 - All nodes should see the same data at the same time
- Availability:
 - Node failures do not prevent survivors from continuing to operate
- Partition-tolerance:
 - The system continues to operate despite network partitions
- A distributed system can satisfy any two of these guarantees at the same time **but not all three**




CAP Theorem



CAP Theorem

- A simple example:


Hotel Booking: are we double-booking the same room?



CAP Theorem

- A simple example:

Hotel Booking: are we double-booking the same room?



CAP Theorem: Proof

- 2002: Proven by research conducted by Nancy Lynch and Seth Gilbert at MIT




Gilbert, Seth, and Nancy Lynch. "Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services." ACM SIGACT News 33.2 (2002): 51-59.




CAP Theorem: Proof

- A simple proof using two nodes:




CAP Theorem: Proof

- A simple proof using two nodes:



CAP Theorem: Proof

- A simple proof using two nodes:



A gets updated from B





Why this is important?

- The future of databases is **distributed** (Big Data Trend, etc.)
- CAP theorem describes the **trade-offs** involved in distributed systems
- A proper understanding of CAP theorem is essential to **making decisions** about the future of distributed database **design**
- Misunderstanding can lead to **erroneous or inappropriate** design choices






Problem for Relational Database to Scale

- The Relational Database is built on the principle of **ACID** (Atomicity, Consistency, Isolation, Durability)
- It implies that a truly distributed relational database should have **availability, consistency and partition tolerance**.
- Which unfortunately is **impossible** ...



Revisit CAP Theorem

- Of the following three guarantees potentially offered by distributed systems:
 - Consistency
 - Availability
 - Partition tolerance
- Pick two
- This suggests there are three kinds of distributed systems:
 - CP
 - AP
 - CA




CAP Theorem 12 year later

- Prof. Eric Brewer: father of CAP theorem
 - “The “2 of 3” formulation was always **misleading** because it tended to oversimplify the tensions among properties. ...
 - **CAP prohibits only a tiny part of the design space:** *perfect availability and consistency in the presence of partitions, which are rare.*”



Consistency or Availability

- Consistency and Availability is not “binary” decision
- AP systems relax consistency in favor of availability
 - but are not inconsistent
- CP systems sacrifice availability for consistency- but are not unavailable
- This suggests both AP and CP systems can offer a degree of consistency, and availability, as well as partition tolerance





AP: Best Effort Consistency

- Example:
 - Web Caching
 - DNS
- Trait:
 - Optimistic
 - Expiration/Time-to-live
 - Conflict resolution



CP: Best Effort Availability

- Example:
 - Majority protocols
 - Paxos, Raft
 - Distributed Locking
 - Google's Chubby: Coarse-grained lock service for loosely-coupled systems
- Feature:
 - Pessimistic locking
 - Make minority partition unavailable
 - E.g.) elect a new leader if the leader is not reachable from majority



Types of Consistency

- **Strong Consistency**
 - After the update completes, **any subsequent access** will return the **same** updated value.
- **Weak Consistency**
 - It is **not guaranteed** that subsequent accesses will return the updated value.
- **Eventual Consistency**
 - Specific form of weak consistency
 - It is guaranteed that if **no new updates** are made to object, **eventually** all accesses will return the last updated value (e.g., *propagate updates to replicas in a lazy fashion*)



Eventual Consistency - A Facebook Example


- Bob finds an interesting story and shares with Alice by posting on her Facebook wall
- Bob asks Alice to check it out
- Alice logs in her account, checks her Facebook wall but finds nothing is there.



Wall
facebook.



?



Eventual Consistency - A Facebook Example

- Bob tells Alice to wait a bit and check out later
- Alice waits for a minute or so and checks back:
 - She finds the story Bob shared with her!





Eventual Consistency - A Facebook Example

- Reason: it is possible because Facebook, Gmail, and many web services use an **eventual consistent model**
- Why Facebook chooses eventual consistency model over the strong consistent one?
 - Facebook has more than 1 billion active users
 - It is non-trivial to efficiently and reliably store the huge amount of data generated at any given time
 - Eventual consistency model offers the option to **reduce the load and improve availability**



Eventual Consistency - A Dropbox Example

- Dropbox enabled **immediate consistency** via synchronization in many cases.
- However, what happens in case of a network partition?



www.bigstock.com - 30744092



Eventual Consistency - A Dropbox Example

- Dropbox embraces eventual consistency:
 - Immediate consistency is impossible in case of a network partition
 - Users will feel bad if their word documents freeze each time they hit Ctrl+S , simply due to the large latency to update all devices across WAN
 - Dropbox is oriented to **personal syncing**, not on collaboration, so it is not a real limitation.



Eventual Consistency - An ATM Example

- In design of automated teller machine (ATM):
 - Strong consistency appear to be a nature choice
 - However, in practice, A beats C
 - Higher availability means higher revenue
 - ATM will allow you to withdraw money even if the machine is partitioned from the network
 - However, it puts a limit on the amount of withdraw (e.g., \$200)
 - The bank might also charge you a fee when a overdraft happens





What if there are no partitions?

- Tradeoff between **Consistency** and **Latency**:
- Caused by the **possibility of failure** in distributed systems
 - High availability → replicate data → consistency problem
- Basic idea:
 - Availability and latency are arguably **the same thing**: unavailable → extreme high latency
 - Achieving different levels of consistency/availability takes different amount of time




CAP → PACELC Theorem

- A more complete description of the space of potential tradeoffs for distributed system:
 - If there is a **partition (P)**,
trade off **availability and consistency (A and C)**;
else (E)
(i.e., when the system is running normally in the absence of partitions,) trade off **latency (L) and consistency (C)**


Abadi, Daniel J. "Consistency tradeoffs in modern distributed database system design." Computer-IEEE Computer Magazine 45.2 (2012): 37.



PACELC



Partitioned



Normal