



# Programming Language & Compiler

## Memory Management & Garbage Collection

**Hwansoo Han**

# Good Memory Management

---

## ▶ Primary goals

- ▶ Good time performance for `malloc` and `free`
  - ▶ Ideally should take constant time (not always possible)
  - ▶ Should certainly not take linear time in the number of blocks
- ▶ Good space utilization
  - ▶ User allocated structures should be large fraction of the heap.
  - ▶ Want to minimize “fragmentation”.


## ▶ Some other goals

- ▶ Good locality properties
  - ▶ Structures allocated close in time should be close in space
  - ▶ “Similar” objects should be allocated close in space
- ▶ Robust
  - ▶ Can check that `free(p1)` is on a valid allocated object `p1`
  - ▶ Can check that memory references are to allocated space

# Internal Fragmentation

---

- ▶ Poor memory utilization caused by *fragmentation*.
  - ▶ Comes in two forms: internal and external fragmentation
- ▶ Internal fragmentation
  - ▶ For some block, internal fragmentation is the difference between the block size and the payload size.



- ▶ Caused by overhead of maintaining heap data structures, padding for alignment purposes, or explicit policy decisions (e.g., not to split the block).
- ▶ Depends only on the pattern of *previous* requests, and thus is easy to measure.

# External Fragmentation

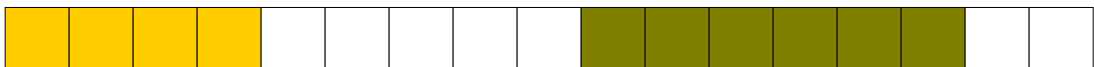
---

- ▶ Occurs when there is enough aggregate heap memory, but no single free block is large enough

`p1 = malloc(4)` 

`p2 = malloc(5)` 

`p3 = malloc(6)` 

`free(p2)` 

`p4 = malloc(6)` 


**oops!**

External fragmentation depends on the pattern of *future* requests, and thus is difficult to measure.

# Knowing How Much to Free


## ▶ Standard method

- ▶ Keep the length of a block in the word preceding the block.
  - ▶ This word is often called the *header field* or *header*
- ▶ Requires an extra word for every allocated block




# Keeping Track of Free Blocks

- ▶ **Method 1:** *Implicit list* using lengths -- links all blocks



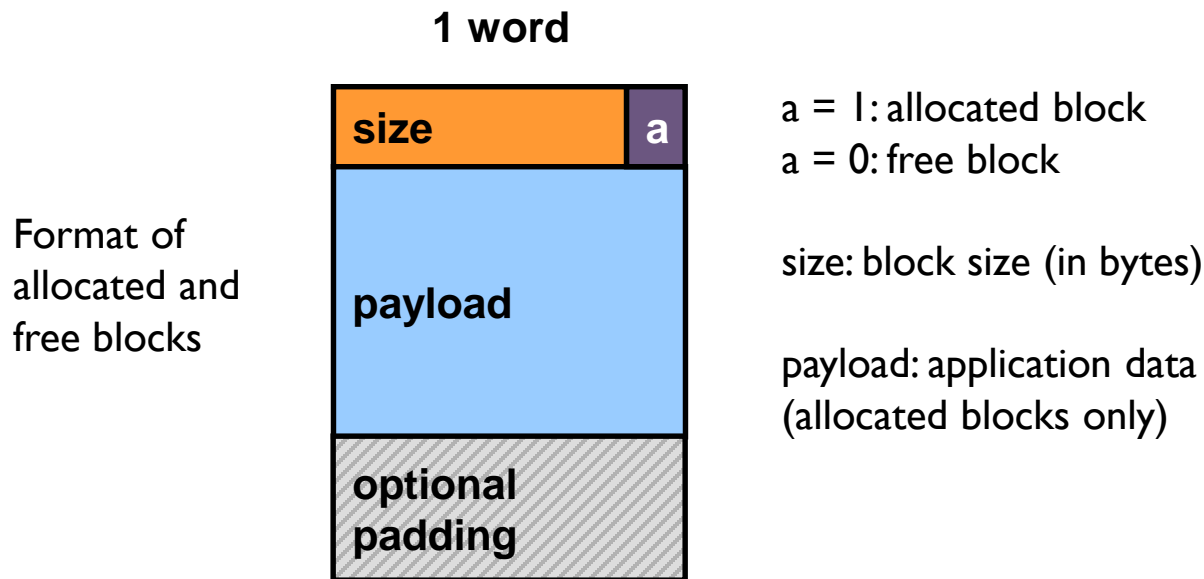
- ▶ **Method 2:** *Explicit list* among the free blocks using pointers within the free blocks



- ▶ **Method 3:** *Segregated free list*
  - ▶ Different free lists for different size classes

# Method 1: Implicit List

- ▶ Need to identify whether each block is free or allocated
  - ▶ Can use extra bit
  - ▶ Bit can be put in the same word as the size if block sizes are always multiples of 2
    - ▶ Mask out low order bit when reading size
    - ▶ If you aligned on 8 bytes, mask out low order 3 bits



# Implicit List: Finding a Free Block

---

## ▶ *First fit:*

- ▶ Search list from beginning, choose first free block that fits
- ▶ Can take linear time in total number of blocks (allocated and free)
- ▶ In practice it can cause “splinters” at beginning of list

## ▶ *Next fit:*

- ▶ Like first-fit, but search list from location of end of previous search

## ▶ *Best fit:*



- ▶ Search the list, choose the free block with the closest size that fits
- ▶ Keeps fragments small --- usually helps avoid fragmentation
- ▶ Will typically run slower than first-fit




# Implicit List: Allocating & Freeing

- ▶ Allocating in a free block – *splitting if needed*

`P = malloc(3)`



`free(p)`




`malloc(5)`

Oops!

# Implicit List: Coalescing

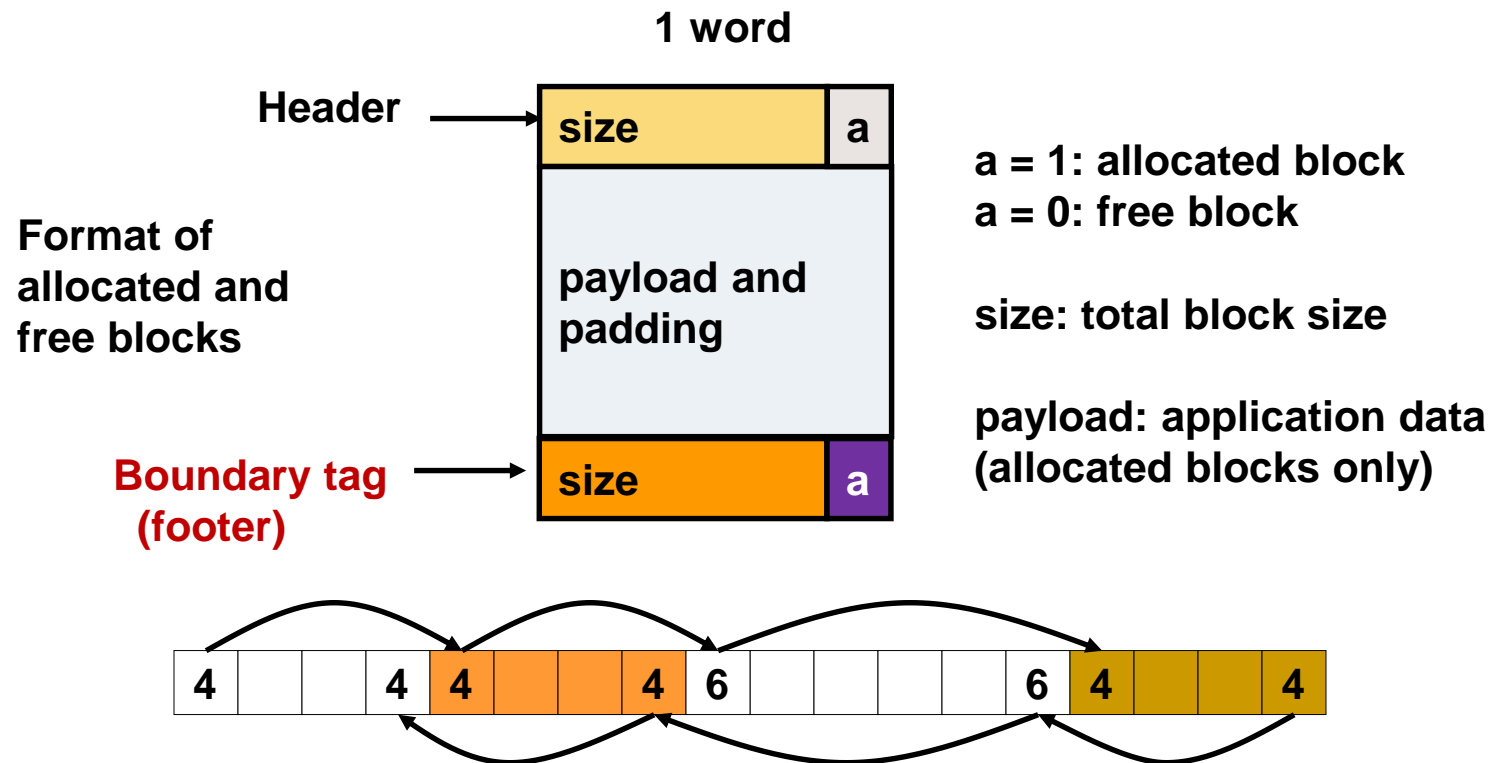
- ▶ Join (*coalesce*) with next and/or previous block if they are free



- ▶ Coalescing with previous free block?

# Implicit List: Bidirectional Coalescing


- ▶ *Boundary tags* [Knuth73]
  - ▶ Replicate size/allocated word at bottom of free blocks
  - ▶ Allows us to traverse the “list” backwards, but requires extra space
  - ▶ Important and general technique!




# Keeping Track of Free Blocks

---

- ▶ Method 1: Implicit list using lengths -- links all blocks



- ▶ Method 2: Explicit list among the free blocks using pointers within the free blocks




- ▶ Method 3: Segregated free lists
  - ▶ Different free lists for different size classes

# Explicit Free Lists





- ▶ Use data space for link pointers
  - ▶ Typically doubly linked
  - ▶ Still need boundary tags for coalescing



- ▶ It is important to realize that links are not necessarily in the same order as the blocks

# Allocating From Explicit Free Lists

---



# Freeing With Explicit Free Lists


---

- ▶ *Insertion policy*: Where in the free list do you put a newly freed block?
  - ▶ LIFO (last-in-first-out) policy
    - ▶ Insert freed block at the beginning of the free list
    - ▶ Pro: simple and constant time
    - ▶ Con: studies suggest fragmentation is worse than address ordered.
  - ▶ Address-ordered policy
    - ▶ Insert freed blocks so that free list blocks are always in address order
      - i.e.  $\text{addr}(\text{pred}) < \text{addr}(\text{curr}) < \text{addr}(\text{succ})$
    - ▶ Con: requires search
    - ▶ Pro: studies suggest fragmentation is better than LIFO


# Keeping Track of Free Blocks

---

- ▶ Method 1: Implicit list using lengths -- links all blocks



- ▶ Method 2: Explicit list among the free blocks using pointers within the free blocks




- ▶ Method 3: Segregated free lists
  - ▶ Different free lists for different size classes



# Segregated Storage

- ▶ Each **size class** has its own collection of blocks



- ▶ General principles
  - ▶ Often have separate size class for every small size (2,3,4,...)
  - ▶ For larger sizes typically have a size class for each power of 2
- ▶ 128 size classes for Doug Lea's malloc.c
  - ▶ 63 exact bins (spaced by 8 byte) : 16, 24, 32, ..., 512
  - ▶ 64 sorted bins (approx. logarithmically spaced) : 576, 640, ...  $2^{31}$

# Segregated Fits

---

- ▶ Array of free lists, each one for some size class
- ▶ To allocate a block of size  $n$ :
  - ▶ Search appropriate free list for block of size  $m > n$
  - ▶ If an appropriate block is found:
    - ▶ Split block and place fragment on appropriate list (optional)
  - ▶ If no block is found, try next larger class
  - ▶ Repeat until block is found
- ▶ To free a block:
  - ▶ Coalesce and place on appropriate list (optional)
- ▶ Tradeoffs
  - ▶ Faster search than sequential fits
  - ▶ Controls fragmentation of simple segregated storage
  - ▶ Coalescing can increase search times
    - ▶ Deferred coalescing can help

# GC - Automated Free for Heap Objects

---

- ▶ **Garbage collection**

- ▶ Automatically reclaim the space that the running program can never access again
- ▶ Performed by the runtime system

- ▶ **Two parts of a garbage collector**

- ▶ Garbage detection
- ▶ Reclamation of the garbage objects' storage

# Liveness in GC

---

- ▶ **A root set**

- ▶ Global variables
- ▶ Local variables in the activation stack
- ▶ Any registers used by active procedures

- ▶ **Live objects**


- ▶ Objects on any directed path of pointers from the roots

# Basic Garbage Collection Techniques

---

- ▶ Assumption – heap objects are self-identifying
- ▶ Basic techniques for GC
  - ▶ Reference counting
  - ▶ Mark-Sweep collection
  - ▶ Mark-Compact collection
  - ▶ Copying collection


# Reference Counting



- Keeping track of how many pointers point to each record

# Reference Counting - problems

- ▶ Cost of reference counting
  - ▶ Too many ref-count increments and decrements
- ▶ Fail to reclaim circular structures
  - ▶ Not always effective



# Mark-Sweep Collection

---

- ▶ **Garbage detection**
  - ▶ Traverse the graph of pointer relationships *and*
  - ▶ Mark all reachable objects
- ▶ **Reclamation**
  - ▶ Sweep unmarked objects
- ▶ **Problems**
  - ▶ Fragmentation
  - ▶ Collection cost - proportional to the size of heap
  - ▶ Poor locality of reference



# Mark-Compact Collection

---

- ▶ The same detection phase by marking
- ▶ Objects are compacted
  - ▶ Moving most of the live objects until all of the live objects are contiguous
- ▶ Pros
  - ▶ No fragmentation problem
  - ▶ Allocation order preserved
- ▶ Cons
  - ▶ Slower than mark-sweep
  - ▶ Need fast compacting algorithms


# Copying Collection

---

- ▶ Move the live objects to a contiguous area
  - ▶ Integrate the traversal of the data and the copying process
- ▶ A simple copying collector
  - ▶ “Stop-and-Copy” using semi-spaces

# “Stop-and-Copy” Using Semi-Spaces

- ▶ Subdivided heap into two contiguous semi-spaces
- ▶ When program demands more unused area of the current semi-space
  - ▶ Stop and copy to reclaim space




Cheney's algorithm:  
breadth-first search to  
traverse the reachable data

# Incremental Tracing Collection

---

- ▶ **Interleave GC with program execution**
  - ▶ Small units of garbage collection interleaved with small units of program execution
  - ▶ Needed for real-time applications
- ▶ **Difficulty**
  - ▶ While the collector is tracing out the graph of reachable objects, the graph may change by the running program
- ▶ **Mark-sweep or copying GC can be made incremental**
  - ▶ Tricolor marking

# Tricolor Marking




- ▶ Live objects' color becomes white, grey, and then black
  - ▶ Collection ends when there are no grey objects
  - ▶ All white objects are garbage
- ▶ Need to coordinate the collector with the “mutator”
  - ▶ Invariant - No black object points to a white object

# Infant Mortality

---

- ▶ Most objects die young
- ▶ Marking entire heap space is time-consuming



# Generational Collection


---

- ▶ Most objects live a very short time, while a small percentage of them live much longer
- ▶ When using copying collection, need to avoid much repeated copying of old objects
  - ▶ Divide the heap into generations
    - ▶ The younger generation is typically several times smaller than the old one
  - ▶ Younger generations are collected more often

# Generational Collection in J2SE

---


- ▶ Arrangement of generations
  - ▶ Young generation (nursery)
    - ▶ Eden + two survivor spaces
  - ▶ Tenured generation
  - ▶ Permanent generation
    - ▶ Code area used by JVM - class and method objects






# Default GC in J2SE

1<sup>st</sup> Minor GC




2<sup>nd</sup> Minor GC




Minor GC: copying

3<sup>rd</sup> Minor GC



Move objects to old space when they become tenured

- New Object
- Garbage
- Live Object



Full GC:  
Mark-compact

# GC Algorithms in J2SE 1.4.1+


## ▶ Young generation

- ▶ Copying collector
- ▶ Parallel collector (2P+)


## ▶ Old generation

- ▶ Mark-compact collector
- ▶ Concurrent Mark-Sweep collector (2P+)


Default Copying Collector




Parallel Collector



Default Mark-compact collector



Concurrent Mark-Sweep collector



# Summary

---

- ▶ **Memory allocation**
  - ▶ First fit, next fit, best fit, segregated fit
- ▶ **Free list management**
  - ▶ Implicit list, explicit list, segregated list
  - ▶ Coalescing with boundary tags
- ▶ **Garbage collection**
  - ▶ Reference counting
  - ▶ Mark-sweep, mark-compact, copying-collection
  - ▶ Stop-the-world vs. incremental GC
  - ▶ Generational GC
  - ▶ Concurrent GC